




UNIVERSITY *of* WEST FLORIDA

# COP4634: Systems & Networks I

*Processes*

- Compiled, executable code
- Stored on disk
- **Passive** entity
- Doesn't **do** anything
- Read from disk
- Written to memory
- Execution begun



Done by loader  
no longer program  
now a process

- Definition: A process (= job) is an instance of a program in memory whose instructions are executed sequentially by a CPU.
- Modern operating system (OS) can execute multiple processes concurrently.
  - no real, only virtual concurrent execution on a single core, single CPU
  - OS switches fast between processes giving each process a chance to run on the CPU
- OS maintains a process control block for each process
  - a data structure maintains information about the process

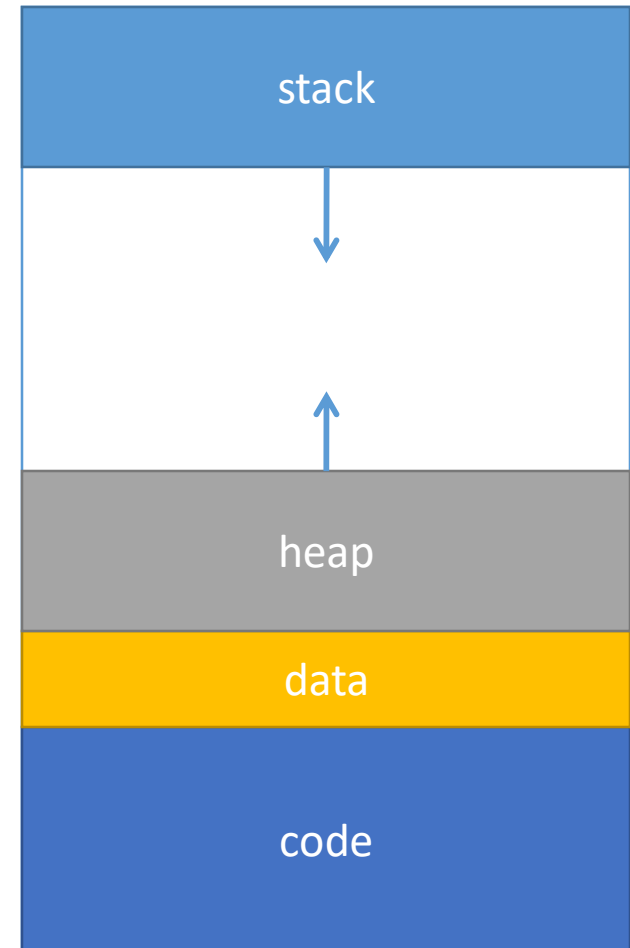
Process consists of

code: a sequence of instructions

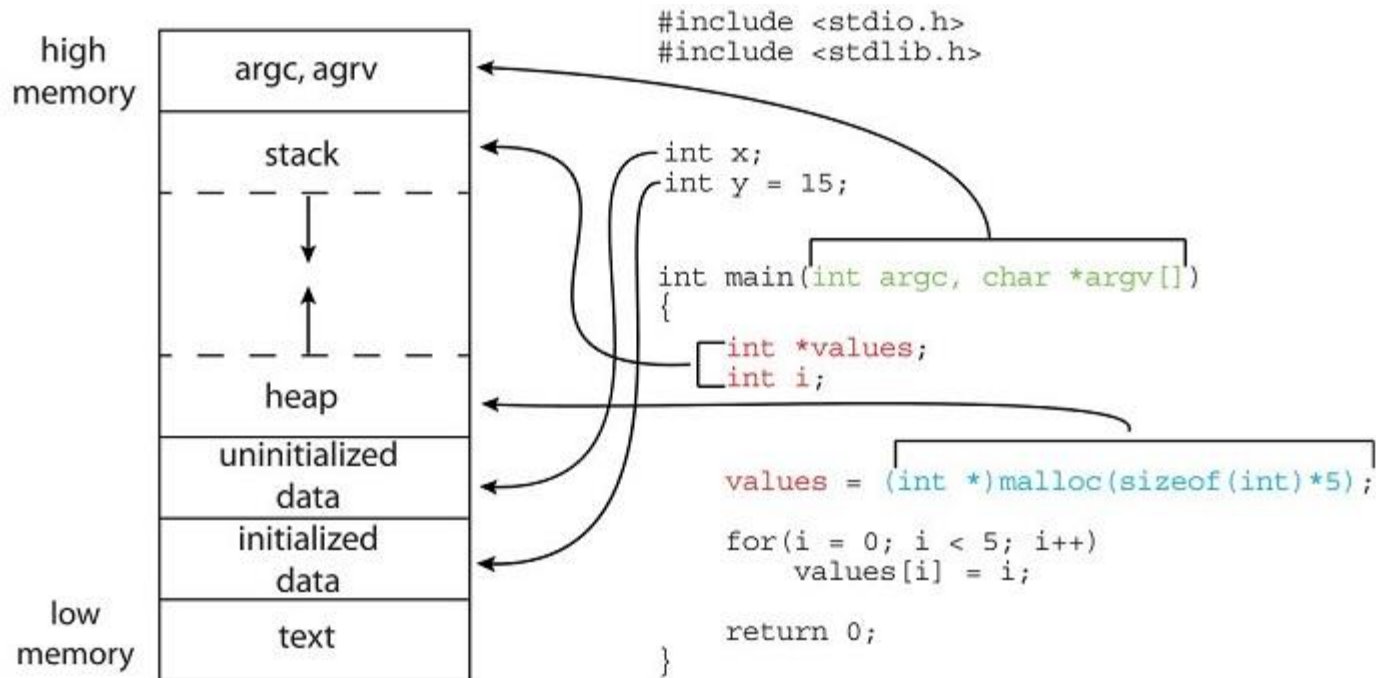
data: global variables

stack: local variables, function  
parameters generated during  
execution

heap: dynamic memory allocated  
during run-time



# Process Layout Illustration



OS maintains a PCB for each process to keep track of a process in memory:

- Process ID,
- Process state (e.g. running, ready, waiting, ...),
- Program counter (address of next instruction),
- CPU registers,
- CPU scheduling information,
- File management (list of open files, working directory, ...),
- I/O status information,
- Memory management information (pointers to text, data, stack).

Process assume different states during execution life cycle:

*New:*

*Running:*

*Waiting:*

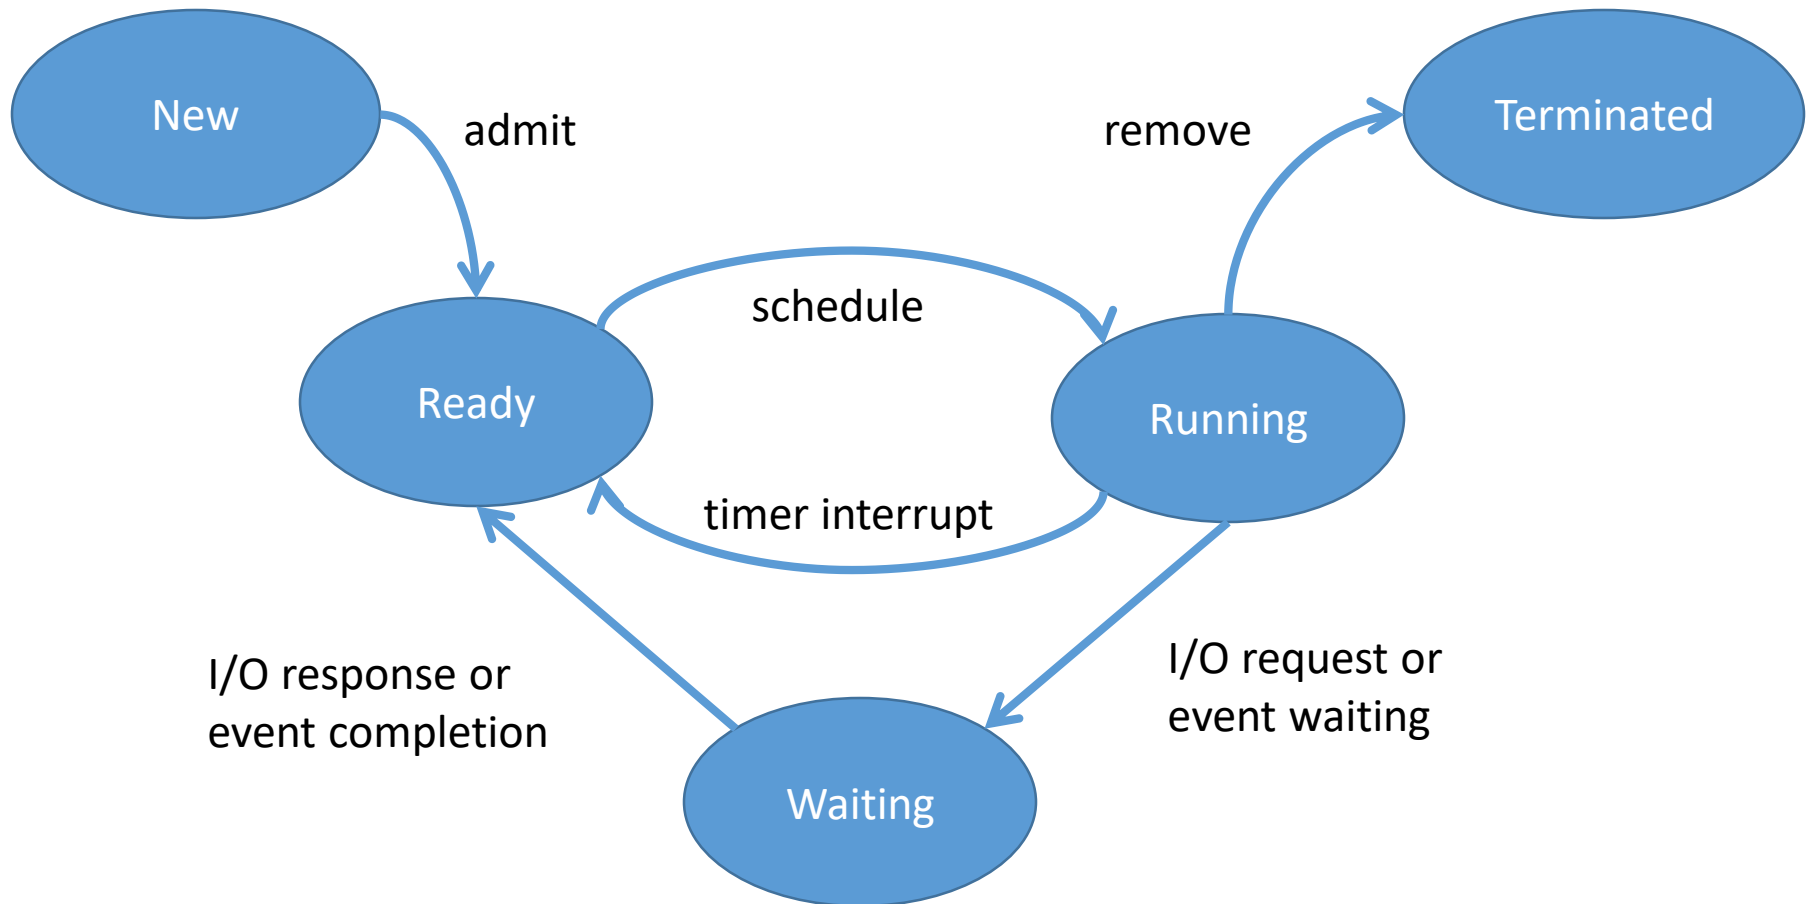
(e.g. an IO device returning some results)

*Ready:*

*Terminated:*

# Process State Diagram (Illustrated)

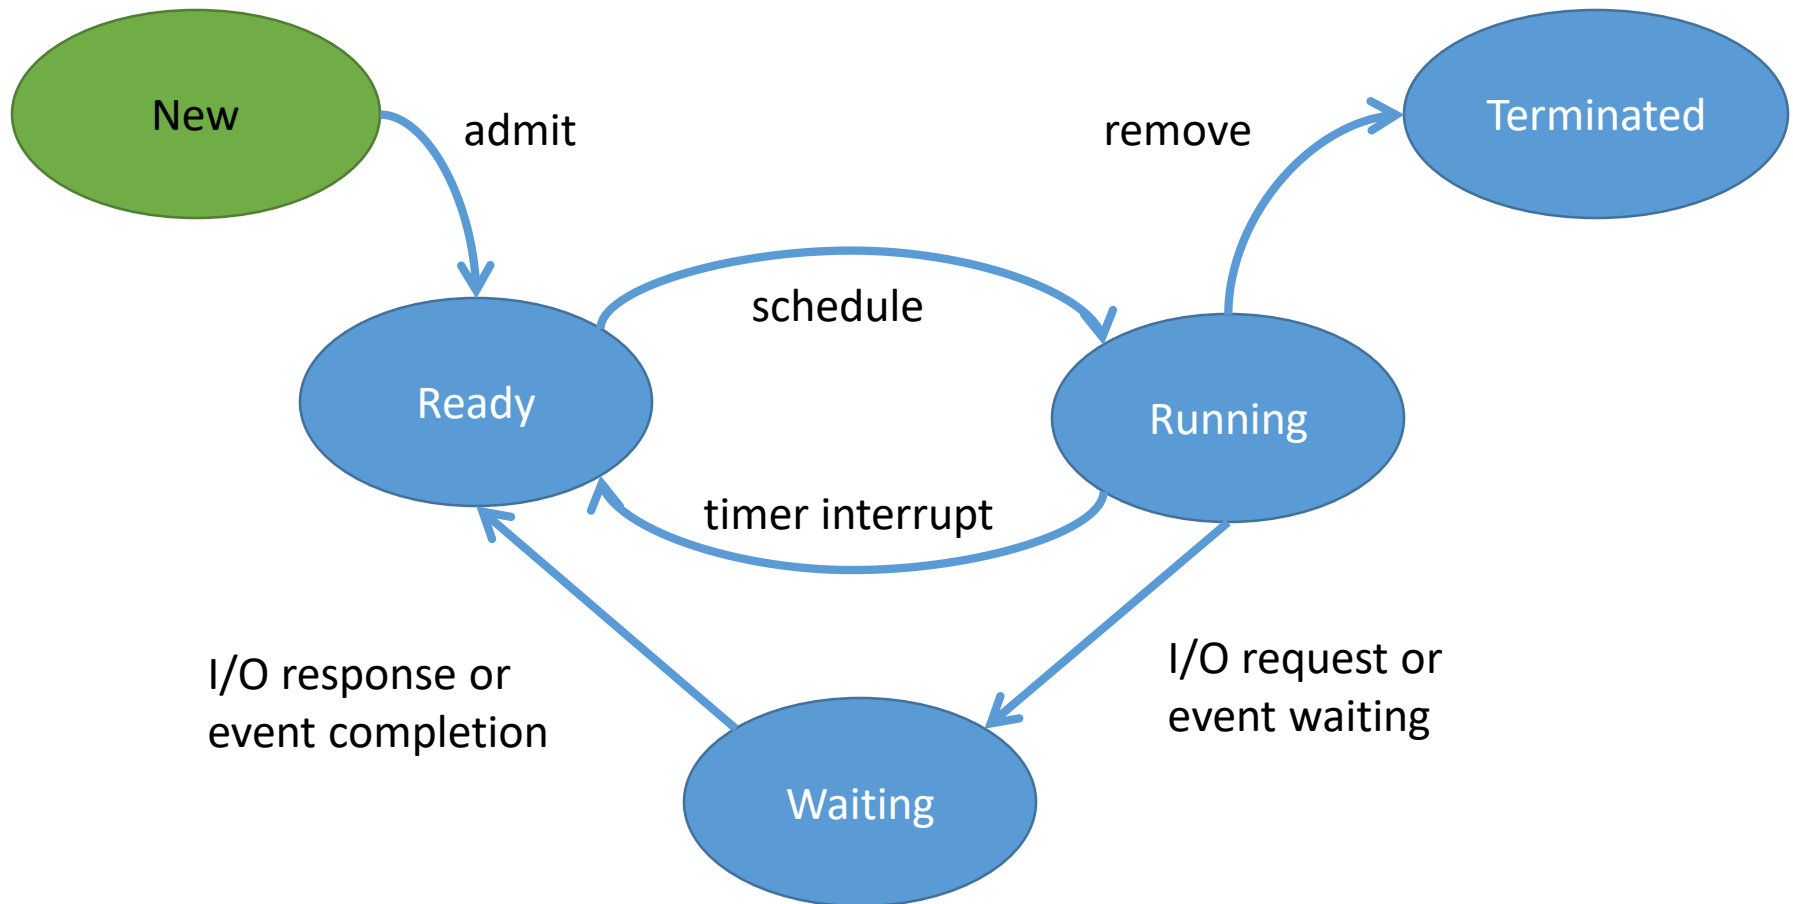
During execution process is in one of five different states.





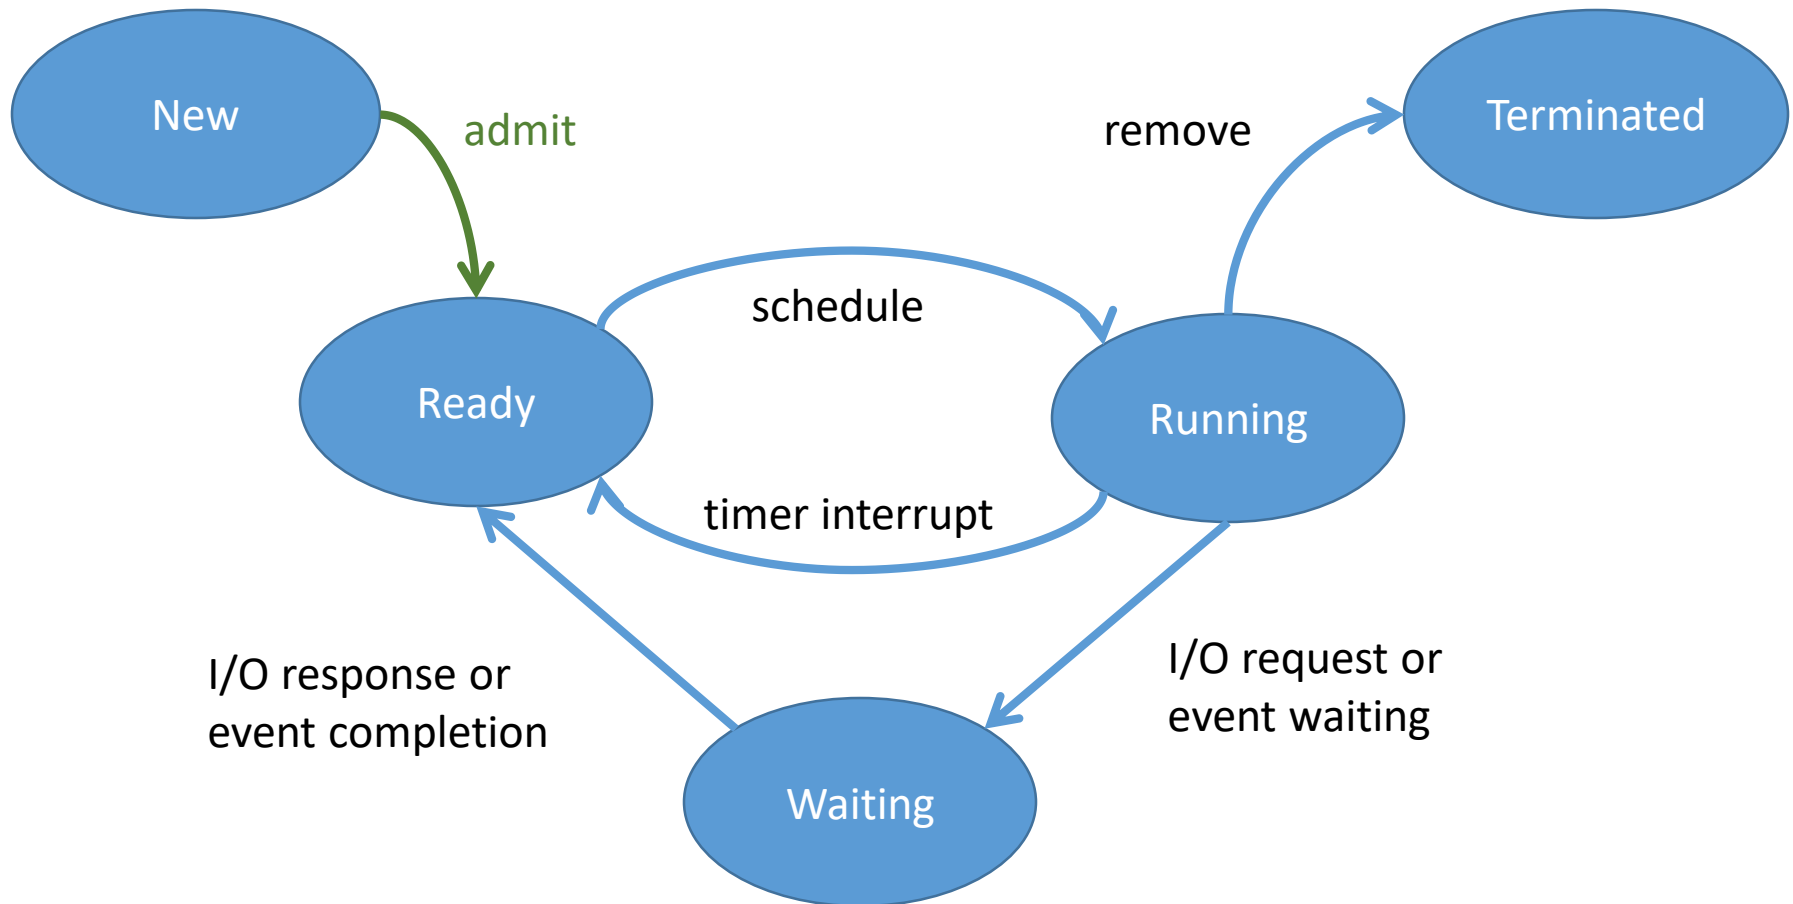
# Process State Diagram (Illustrated)

Kernel is aware of another process requesting to be started.  
No memory has been allocated.



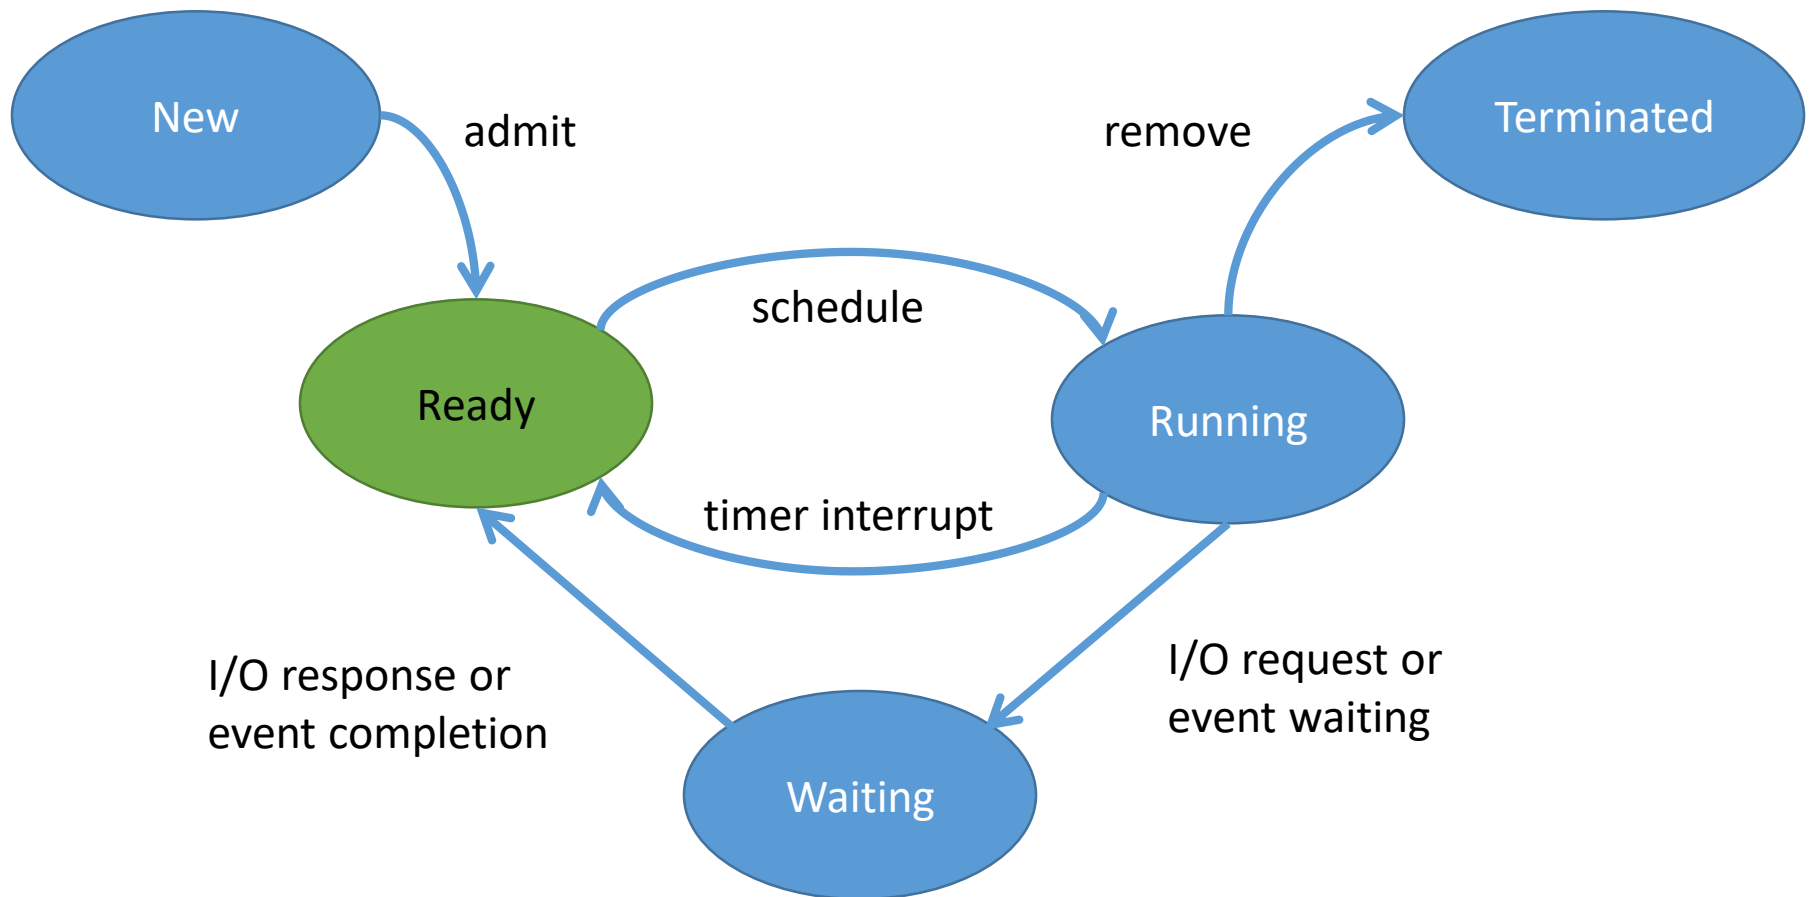
# Process State Diagram (Illustrated)

PCB is allocated and initialized. Memory for the new process is allocated. Process is ready for execution.



# Process State Diagram (Illustrated)

Process is waiting, ready for access to CPU.  
OS CPU scheduler may select process for execution.

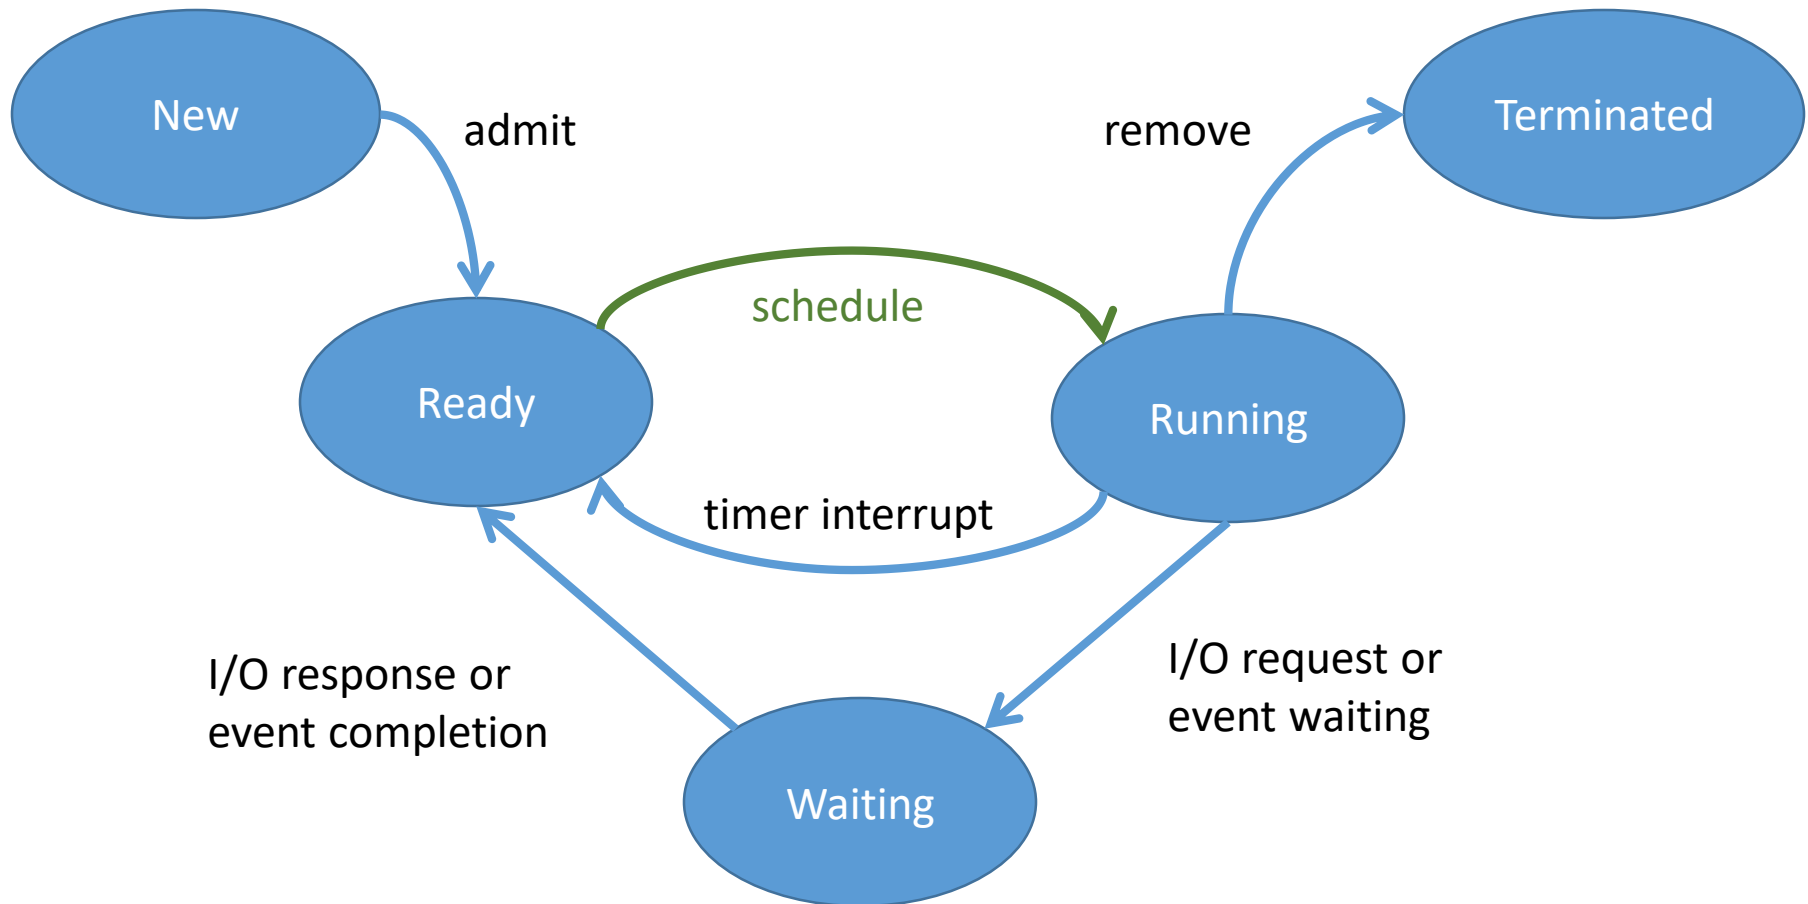


# Process State Diagram (Illustrated)

Process is selected for execution by scheduler.

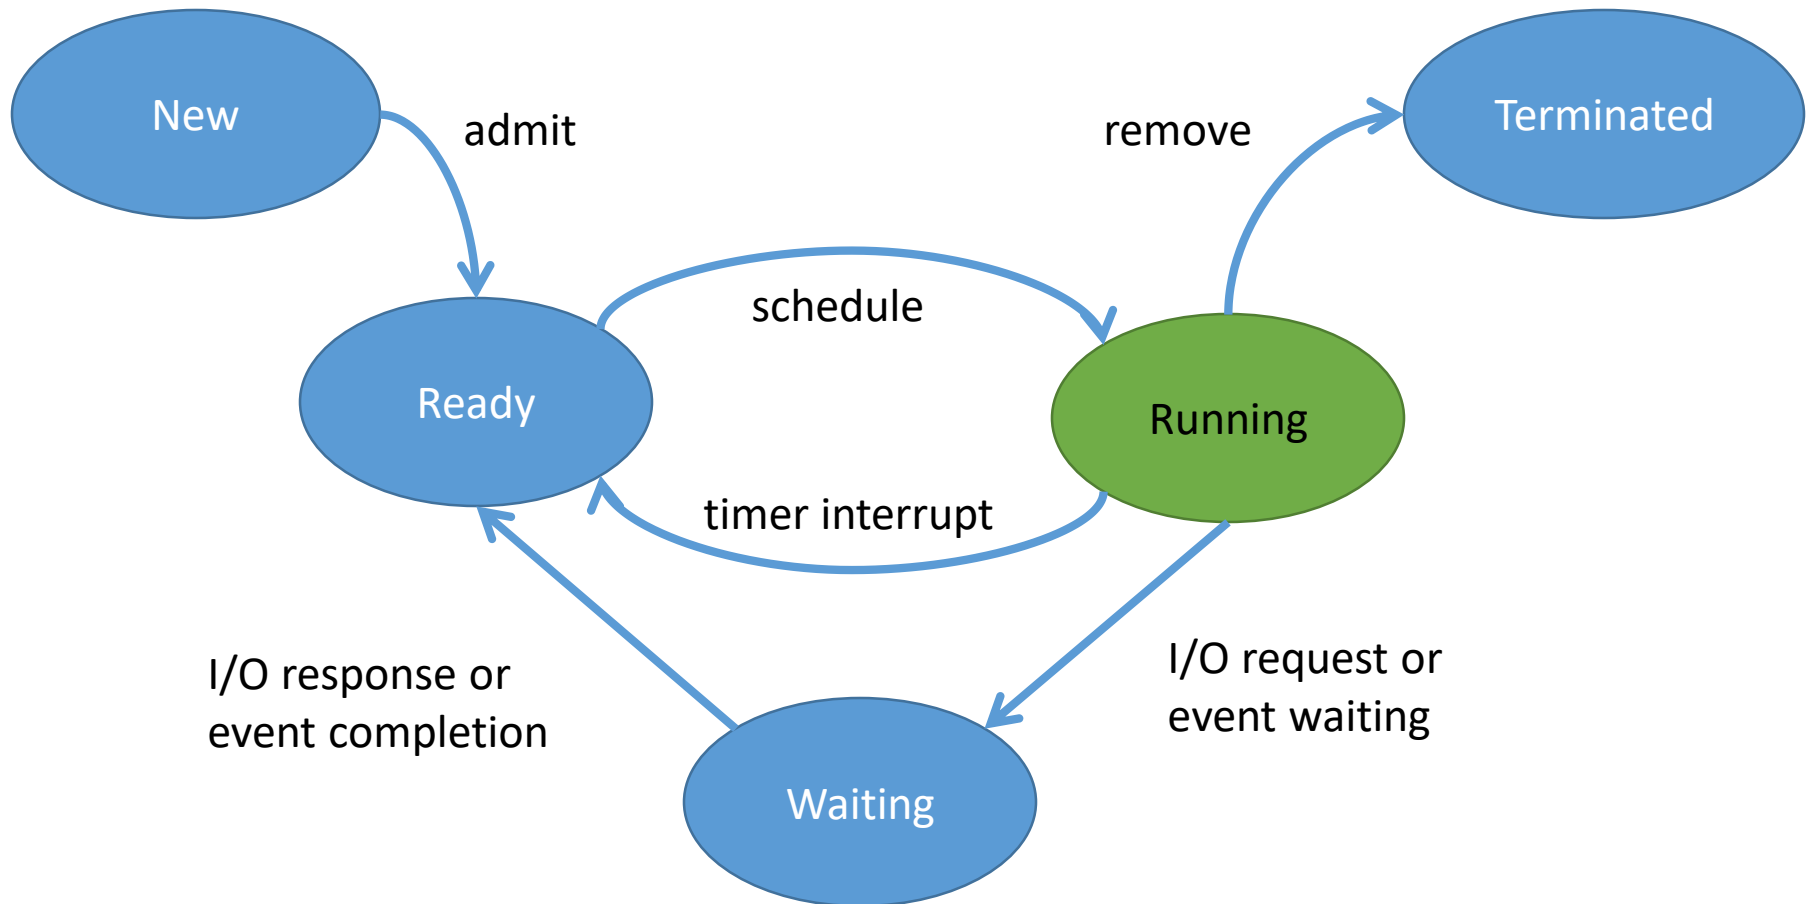
Alarm interrupt is set for *now + quantum*.

Process is “placed” on CPU; state changes to running.



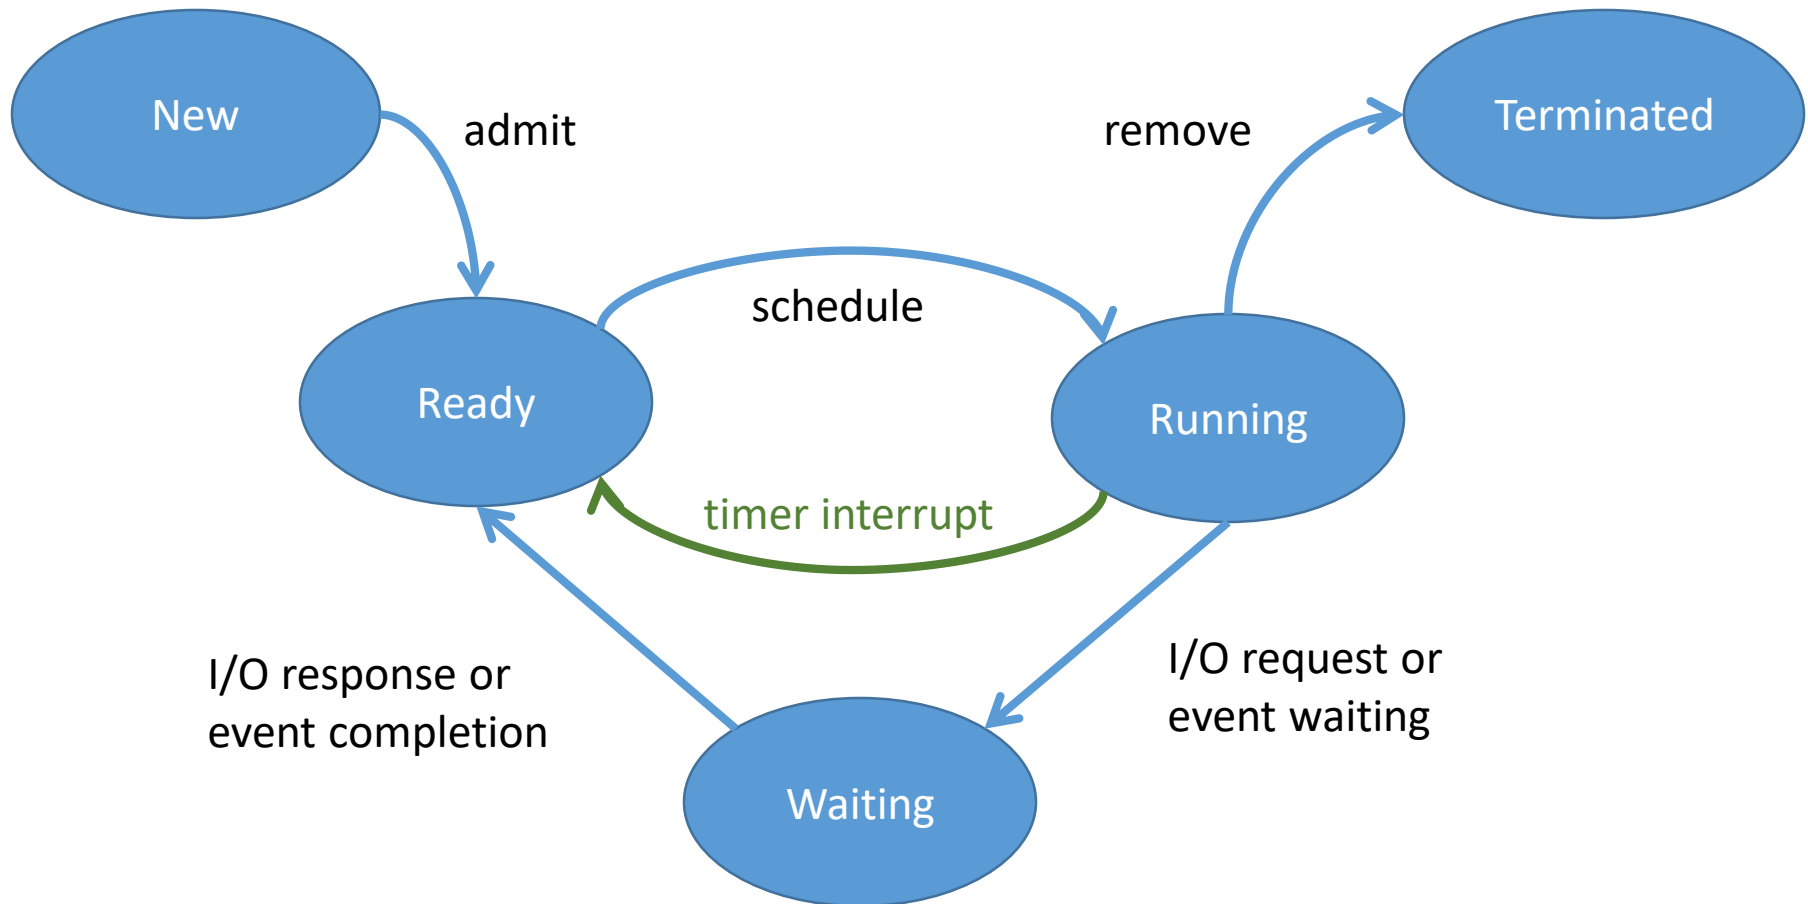
# Process State Diagram (Illustrated)

CPU fetches and executes process instructions. Process may lose CPU access in three different ways.



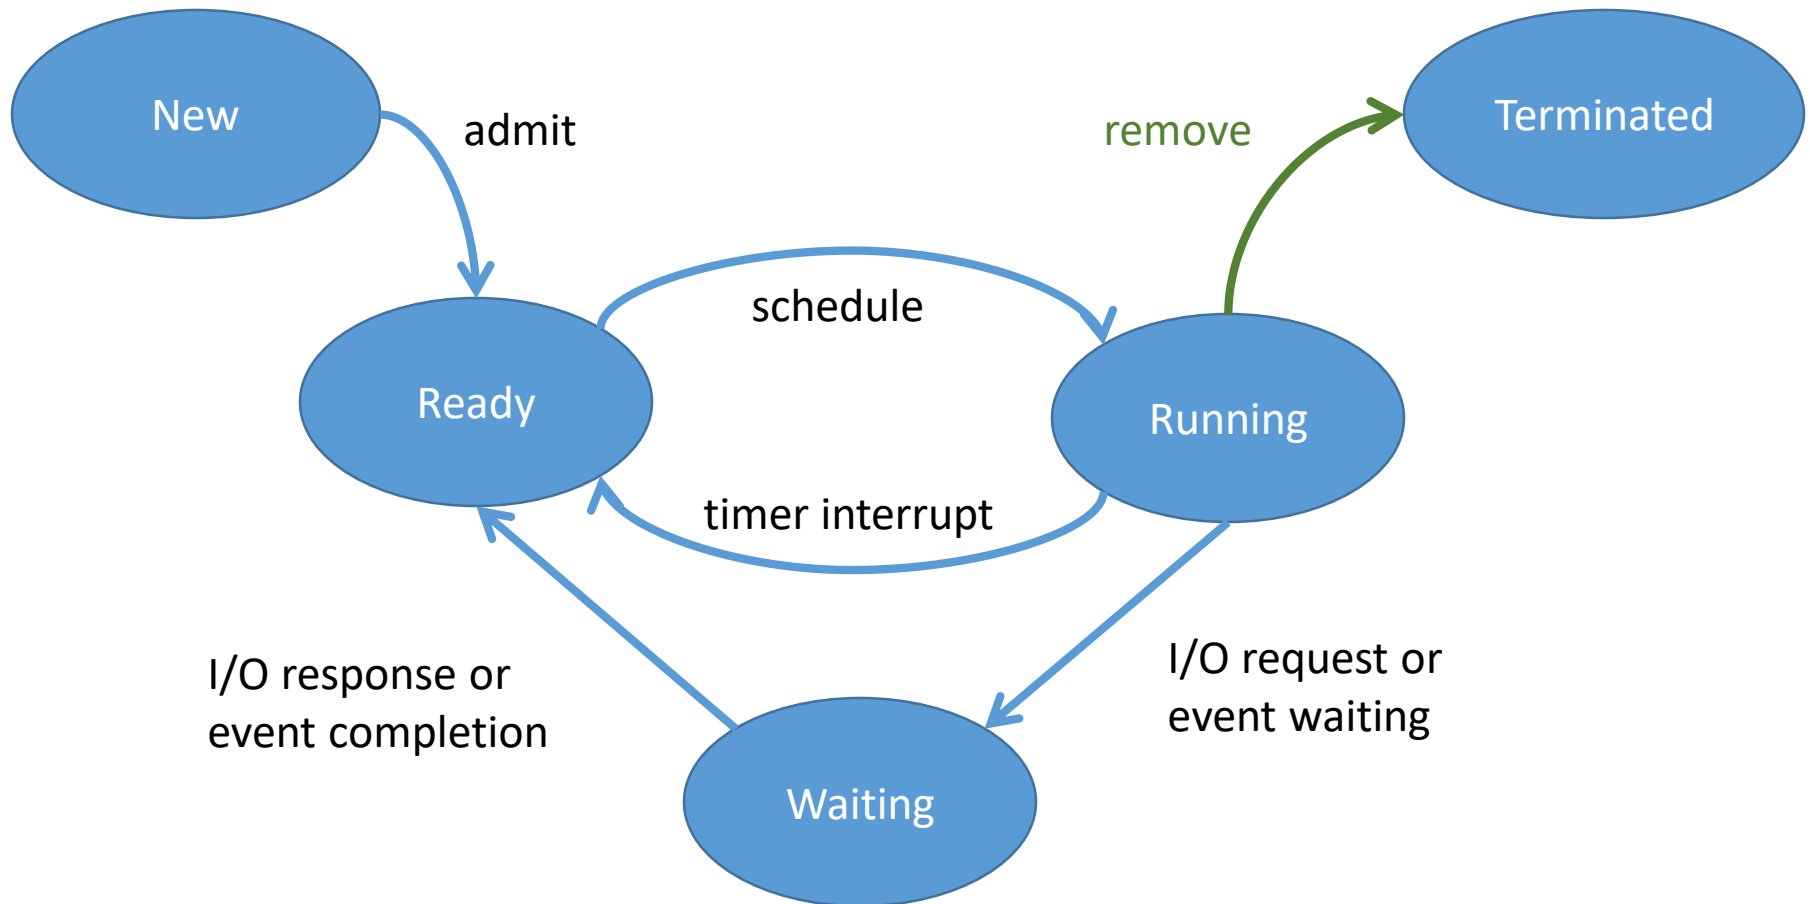
# Process State Diagram (Illustrated)

Quantum expires – Interrupt triggers kernel to resume control.  
Process is moved to *ready list* and next process is *scheduled*.



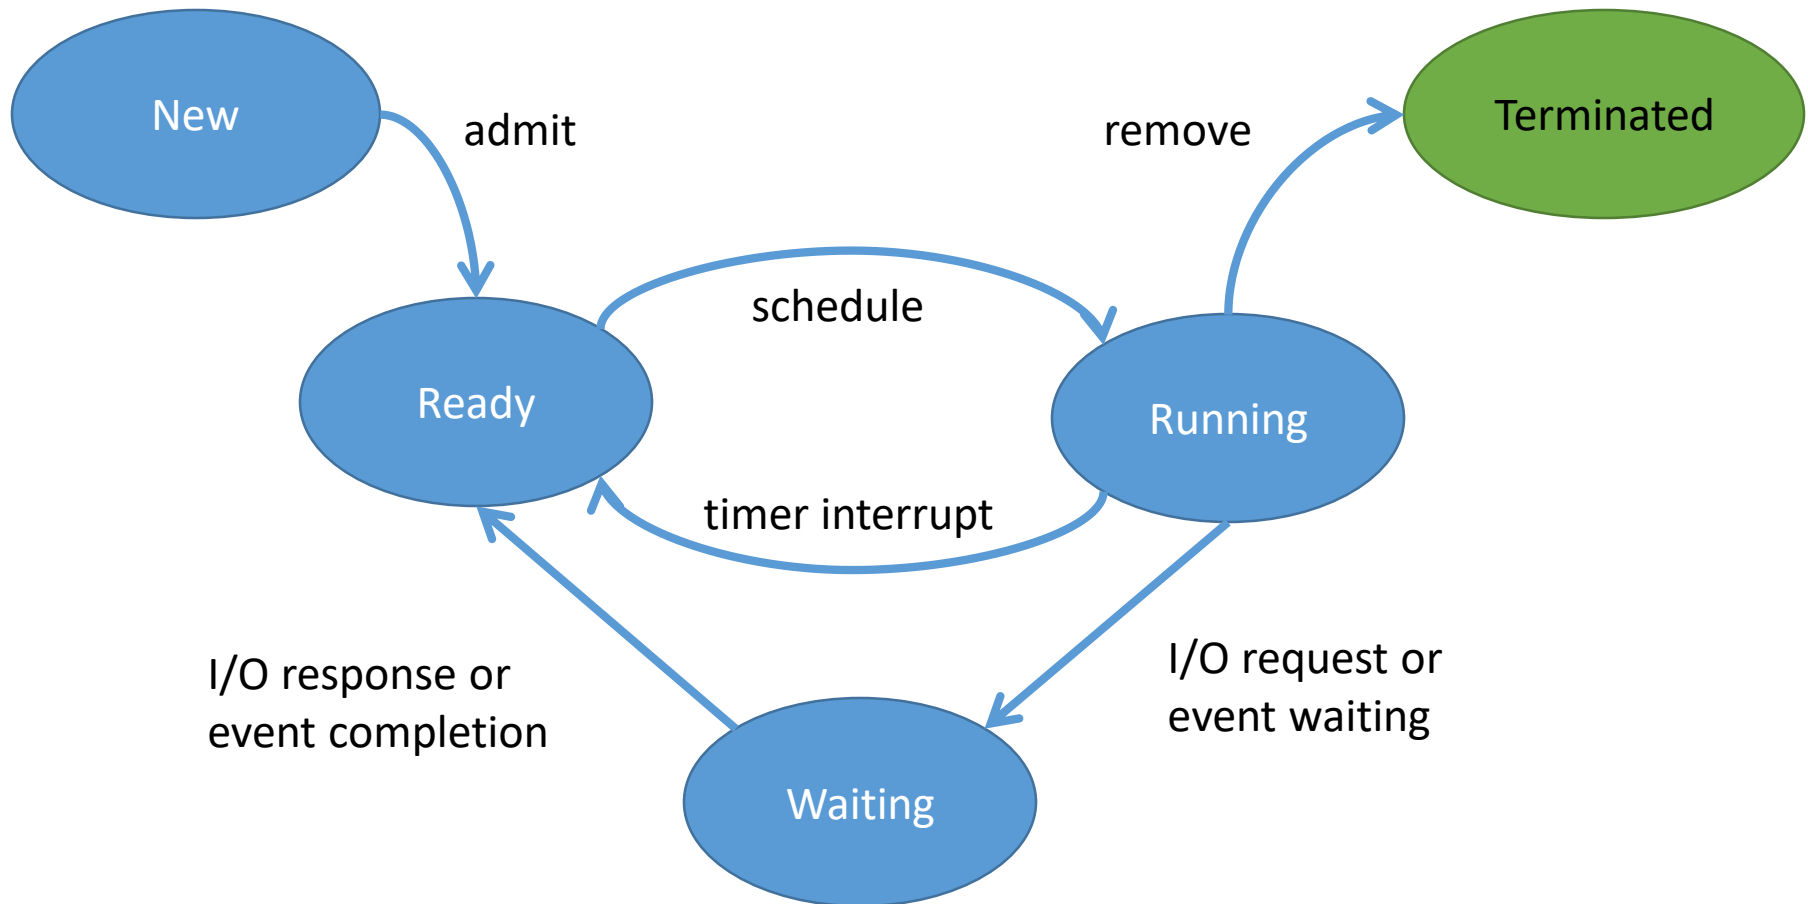
# Process State Diagram (Illustrated)

Process exits – return from `main()`, calls `exit()`, exception.  
Still a process – resources remain allocated until deleted by OS..



# Process State Diagram (Illustrated)

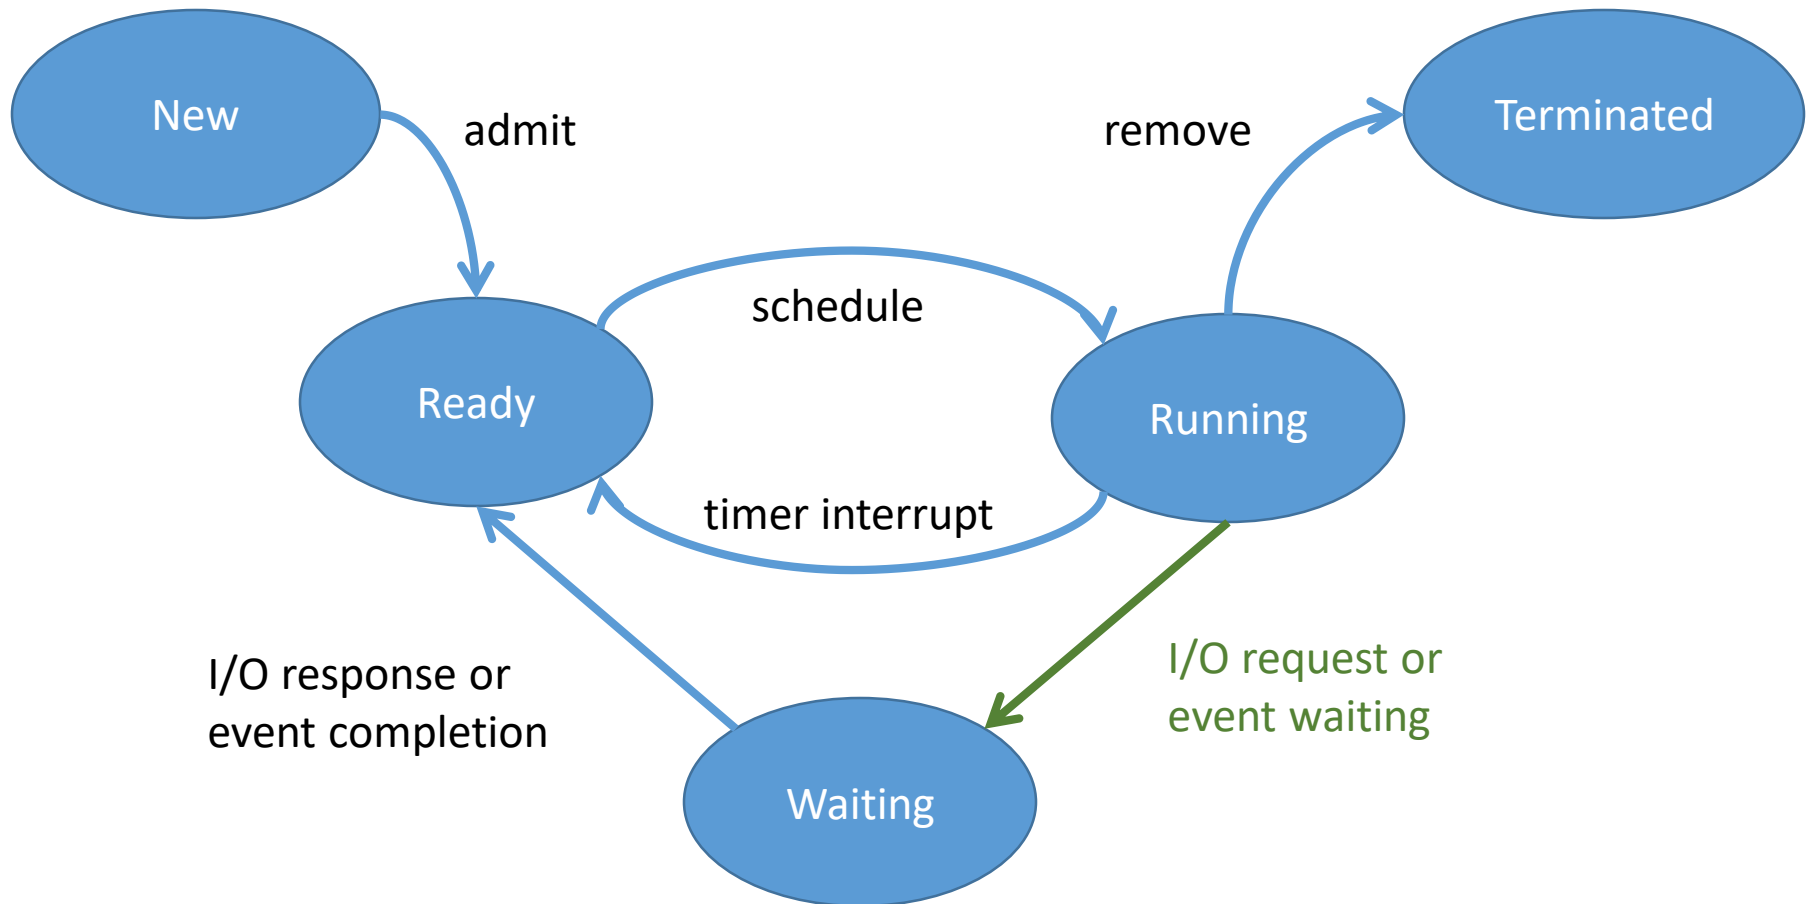
Kernel reclaims resources and return value is given to parent.  
Could become a *zombie* process.





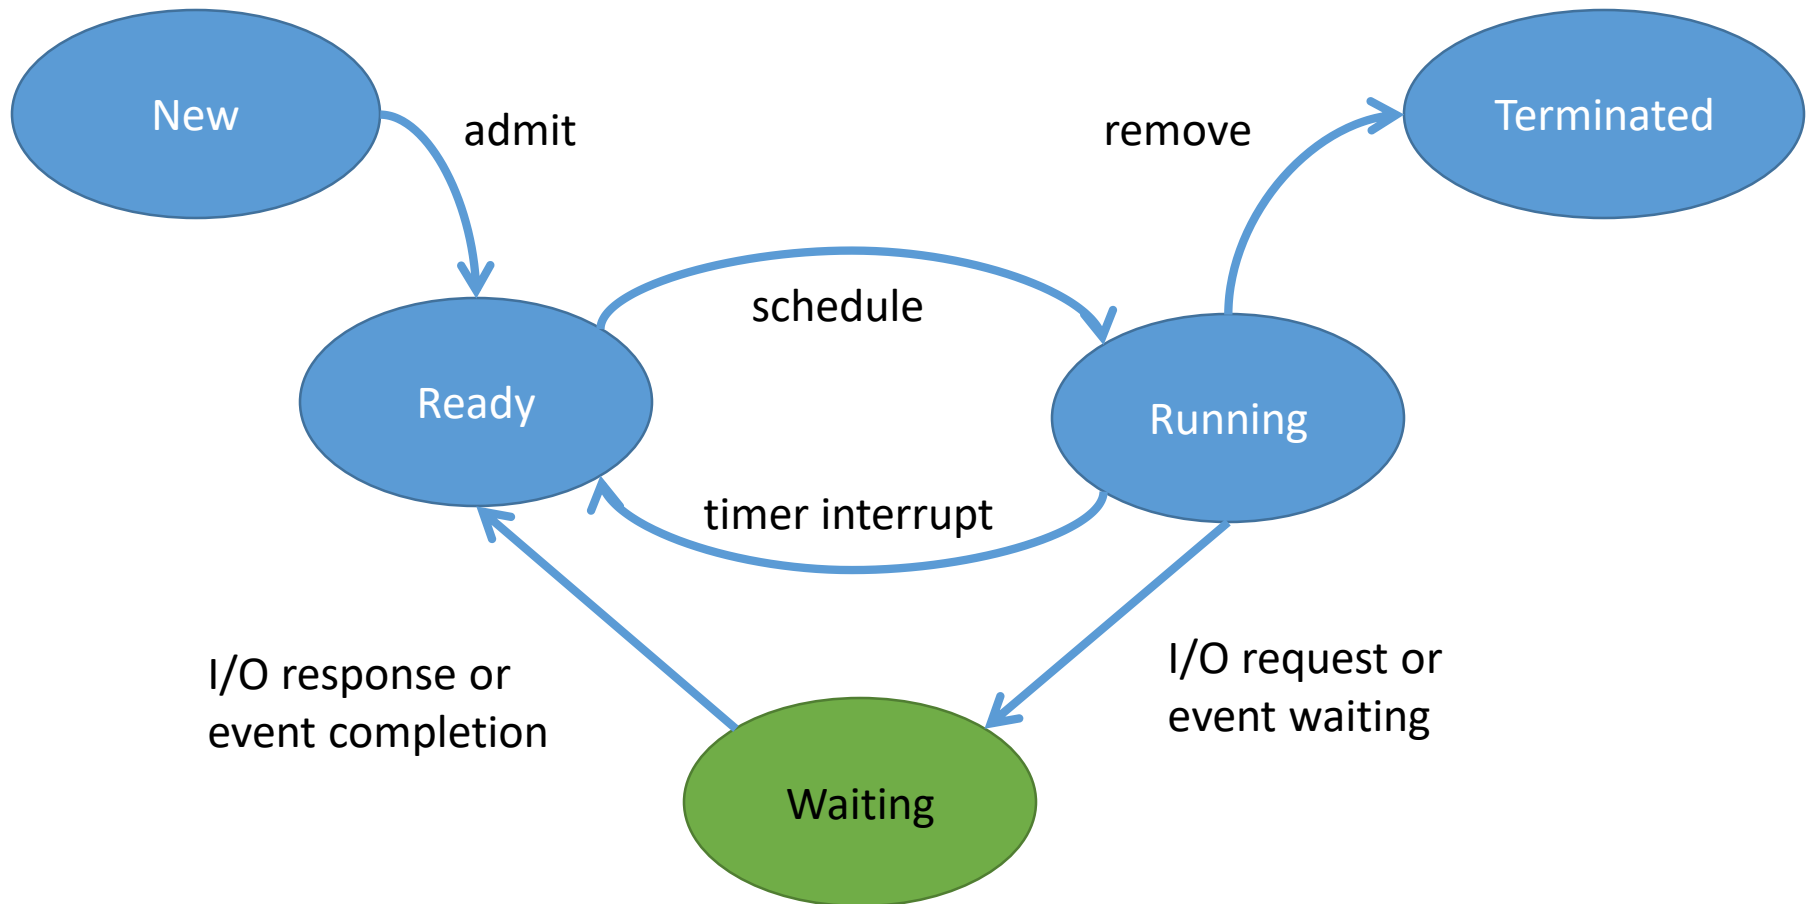
# Process State Diagram (Illustrated)

Process initiates input or output or could *wait()* for other processes.  
Process started something that will take a while.



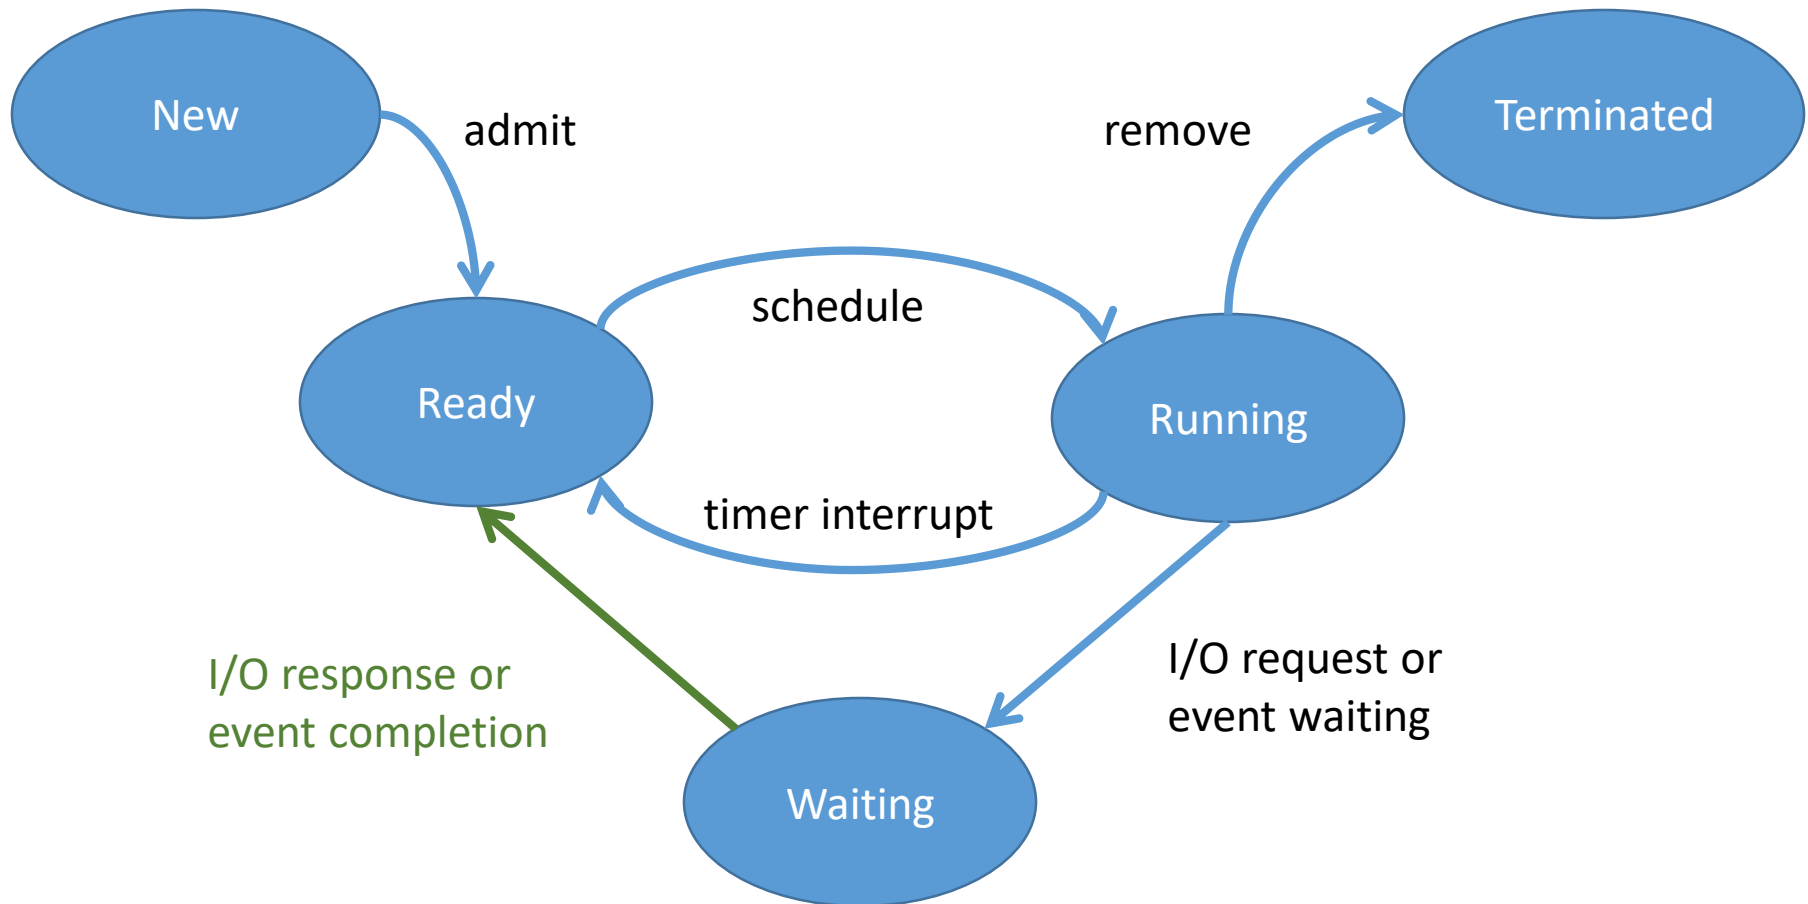
# Process State Diagram (Illustrated)

Process waits for input or output to complete or for another process to wake it up. Process is not ready to continue execution.



# Process State Diagram (Illustrated)

Input/Output or `wait()` is completed; process is ready to be selected for execution.



- General purpose registers (R0, R1, ...)
- Special purpose registers
  - IP/PC: instruction pointer/program counter
  - IR: instruction register
  - SP: stack pointer
- More stuff later

1. Fetch instruction at address PC into IR
2. Increment PC (by 1, 2, 4, ?)
3. Decode instruction in IR
4. Execute instruction in IR
5. *Check for interrupts*
6. Go back to step 1

- SP has address of top of runtime stack
- Runtime stack contains “activation records”
- Stack grows down
- Call to function:
  - creates activation record for function
  - pushes AR onto stack
  - jumps to code for function

## Contains

- parameters
- local variables
- return value
- return address
- other items

```
int add(int x, int y) {  
    int z;  
    z = x + y;  
    return z;  
}
```

```
int main() {  
    int val;  
    val = add(2, 3);  
    printf("%d\n", val);  
    return 0;  
}
```

1. Push arguments onto stack
2. Push local vars onto stack
3. Push return address onto stack
4. Push return value onto stack  
(created activation record)
5. Jump to user-defined code

Limited in what it can do

Only access to user-mode instructions




```
int add(int x, int y) {  
    int z;  
    z = x + y;  
    return z;  
}  
  
int main() {  
    int val;  
    val = add(2, 3);  
    printf("%d\n", val);  
}
```

PC

main()

SP



```
int add(int x, int y) {
    int z;
    z = x + y;
    return z;
}
```

PC

```
int main() {
    int val;
    val = add(2, 3);
    printf("%d\n", val);
}
```

main()

add()

SP

val	?
RV	
RA	
y	3
x	2
z	5
RV	5
RA	

```
int add(int x, int y) {
    int z;
    z = x + y;
    return z;
}
```

PC

```
int main() {
    int val;
    val = add(2, 3);
    printf("%d\n", val);
}
```

main()

add()

SP

val	?
RV	
RA	
y	3
x	2
z	5
RV	5
RA	

```
int add(int x, int y) {  
    int z;  
    z = x + y;  
    return z;  
}  
  
int main() {  
    int val;  
    val = add(2, 3);  
    printf("%d\n", val);  
}
```

PC

main()

SP

val	5
RV	
RA	
	3
	2
	5
	5


User-Mode

add

mult

load

stor



Kernel-Mode

in

out

port

cmod

CPU has execution **mode** bit

- 0 – CPU is currently in kernel-mode
- 1 – CPU is currently in user-mode

`cmod` instruction changes mode  
from kernel to user

No instruction to change  
from user to kernel (Why?)

1. Push arguments onto stack
2. Push local vars onto stack
3. Push return address onto stack
4. Push return value onto stack  
(created activation record)
5. Switch to kernel mode – How?
6. Jump to system-defined code

Access to all instructions

## 1. Software

- System call
- Called a “trap” into the kernel
- Jump to well-known function
- Hardware switches to kernel mode
- Example: call to `printf()` or `read()`



## 2. Exception

- Error state or debugging
- Similar to system call without return (error)
- Jump to well-known function
- Hardware switches to kernel mode
- Example: division by zero or segmentation violation
- Result: core dump

## 3. Hardware

- Called an “interrupt”
- Communication between kernel & devices
- Can occur between any two instructions
- Similar to system call without call
- Jump to well-known function
- Hardware switches to kernel mode
- Example: clock tick or I/O complete

```
int add(int a, int b) {
```

```
    int c;
```

```
    c = a + b;
```

```
    ...
```

```
    LOAD R1, b
```

```
    LOAD R2, a
```

```
    ADD R3, R2, R1
```

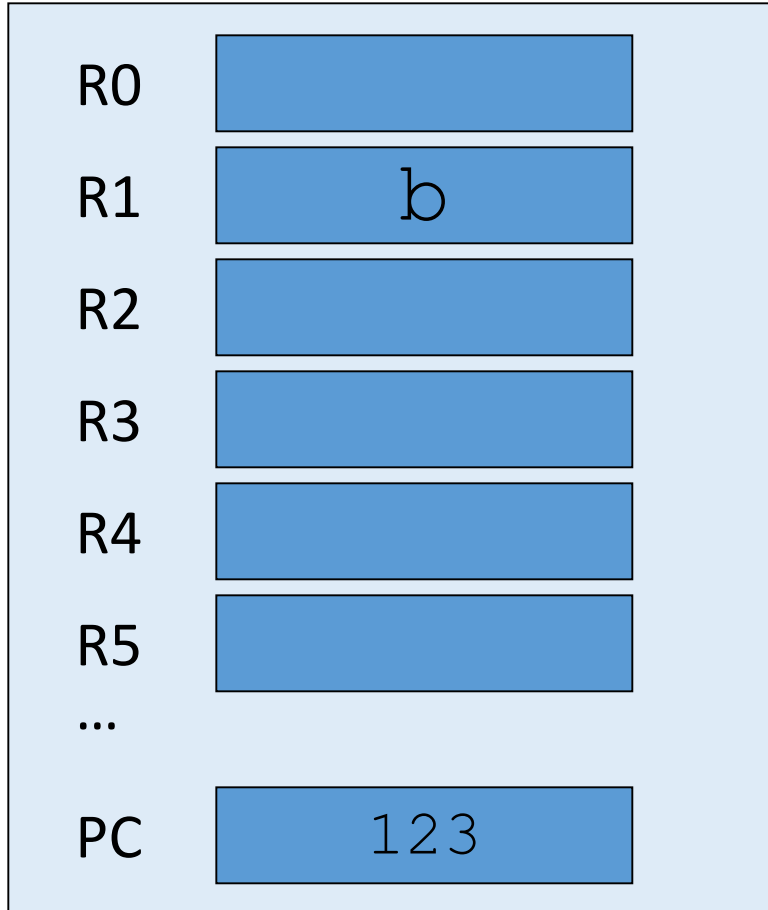
```
    STOR R3, c
```

```
    return c;
```

```
}
```

```
...
```

## CPU



...

LOAD R1, b

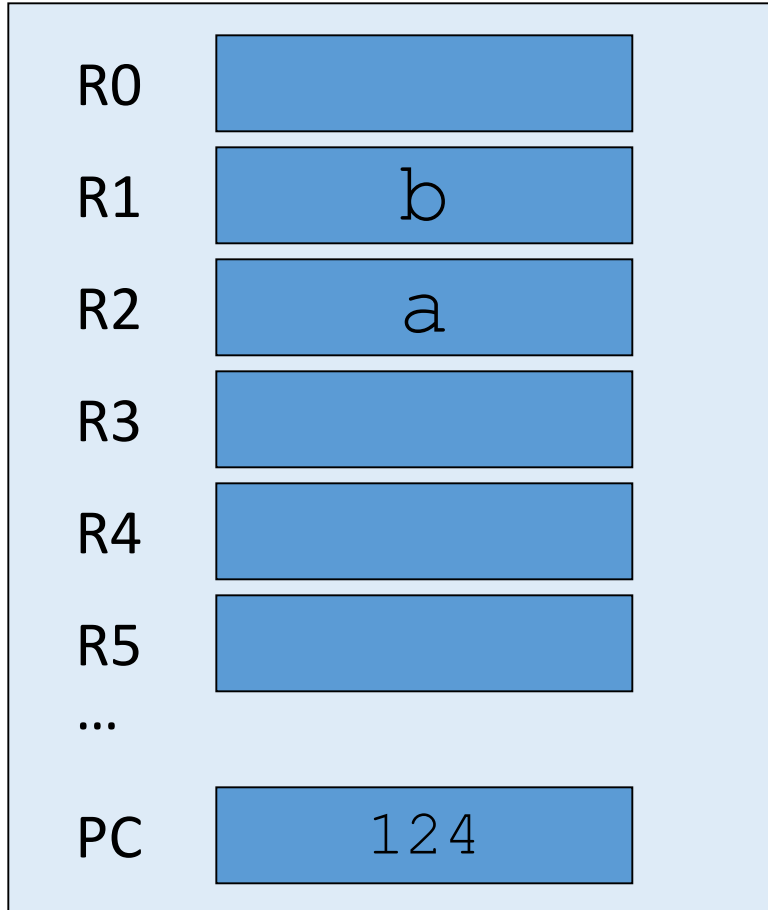
LOAD R2, a

ADD R3, R2, R1

STOR R3, c

...

## CPU



...

LOAD R1, b

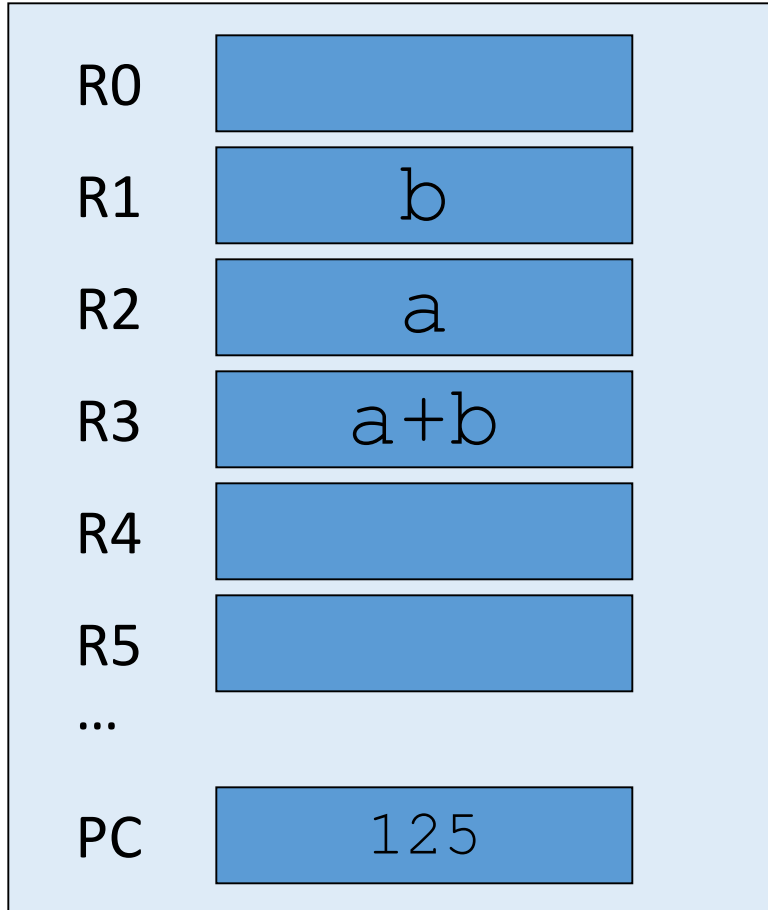
LOAD R2, a

ADD R3, R2, R1

STOR R3, c

...

## CPU



...

LOAD R1, b

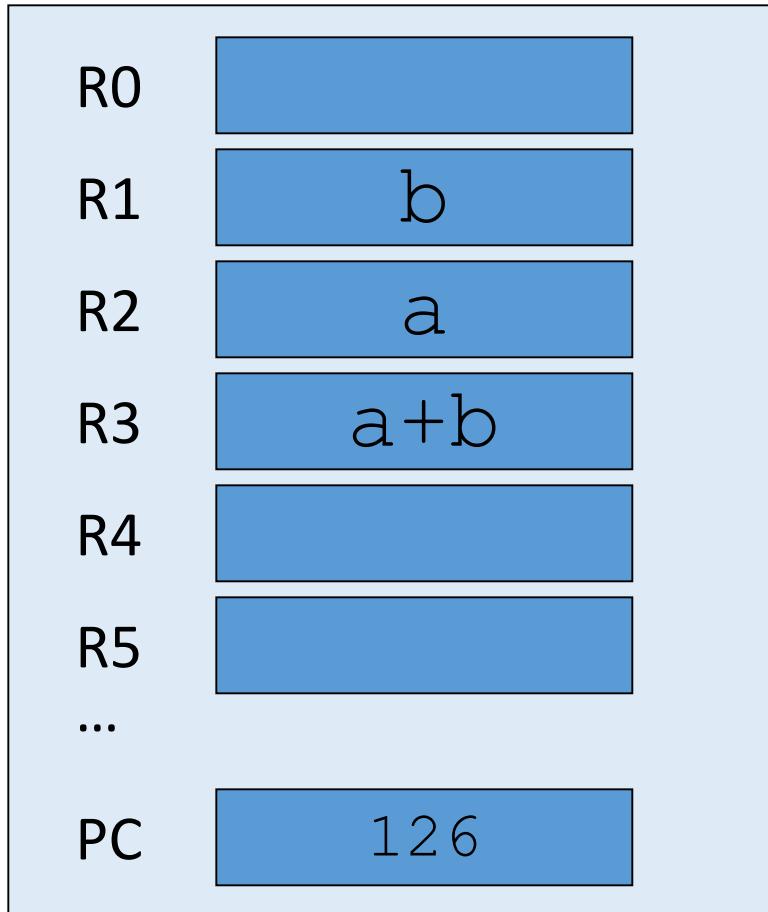
LOAD R2, a

ADD R3, R2, R1

STOR R3, c

...

## CPU



...

LOAD R1, b

LOAD R2, a

ADD R3, R2, R1

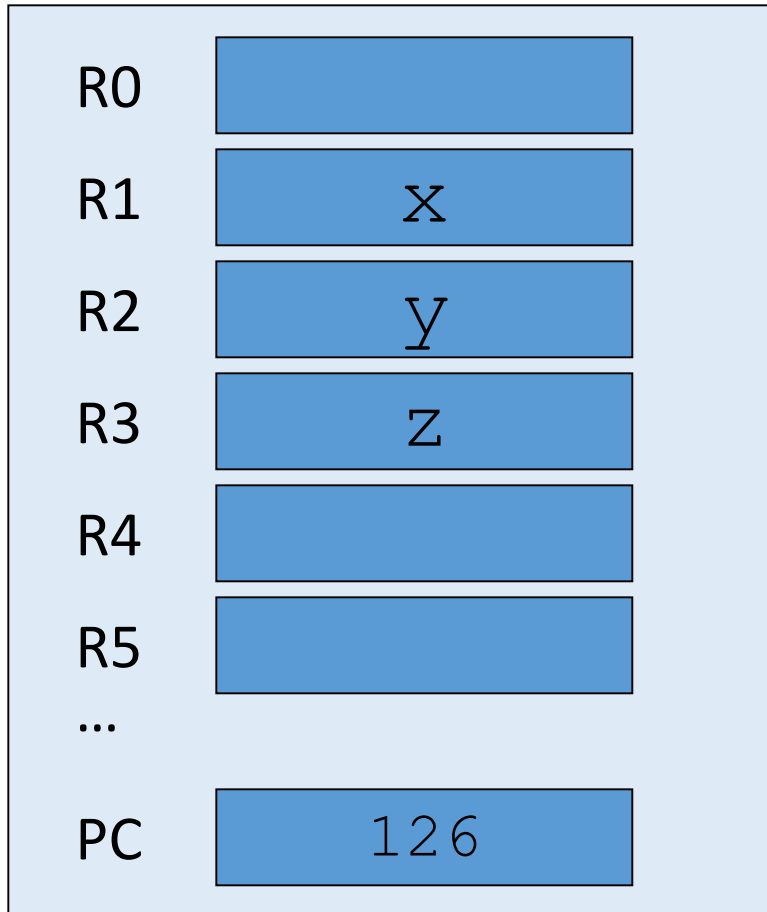
STOR R3, c

...

- PC is 126
- `ADD` has occurred but not `STOR`
- Next scheduled process alters registers
- When our process returns to the CPU...
- this is what it finds



## CPU



...

LOAD R1, b

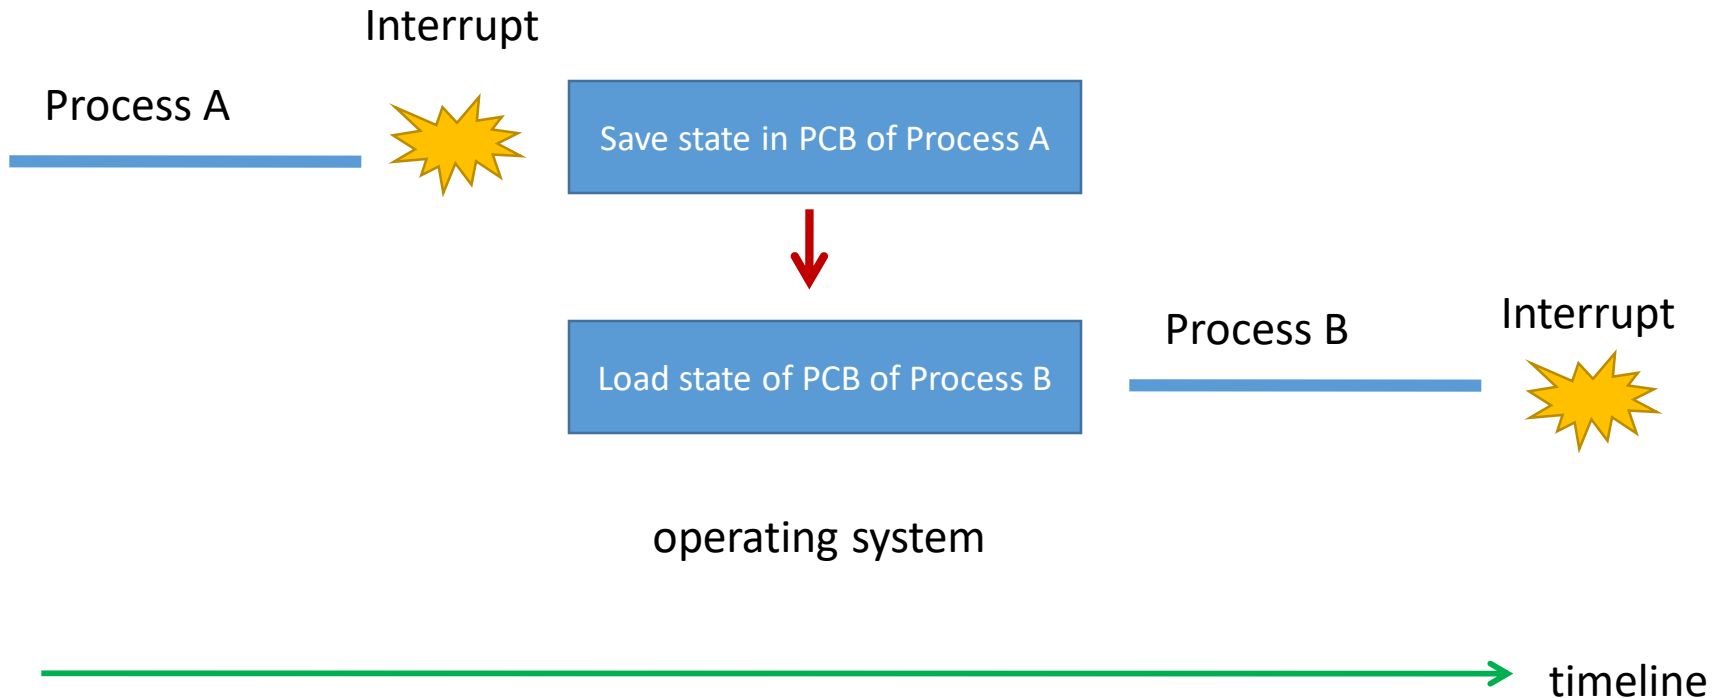
LOAD R2, a

ADD R3, R2, R1

STOR R3, c

...

# CPU Process Switching (Illustration)



Note: switching a process generates management overhead that prevents execution of useful work.

When process  $P_i$  removed from CPU,

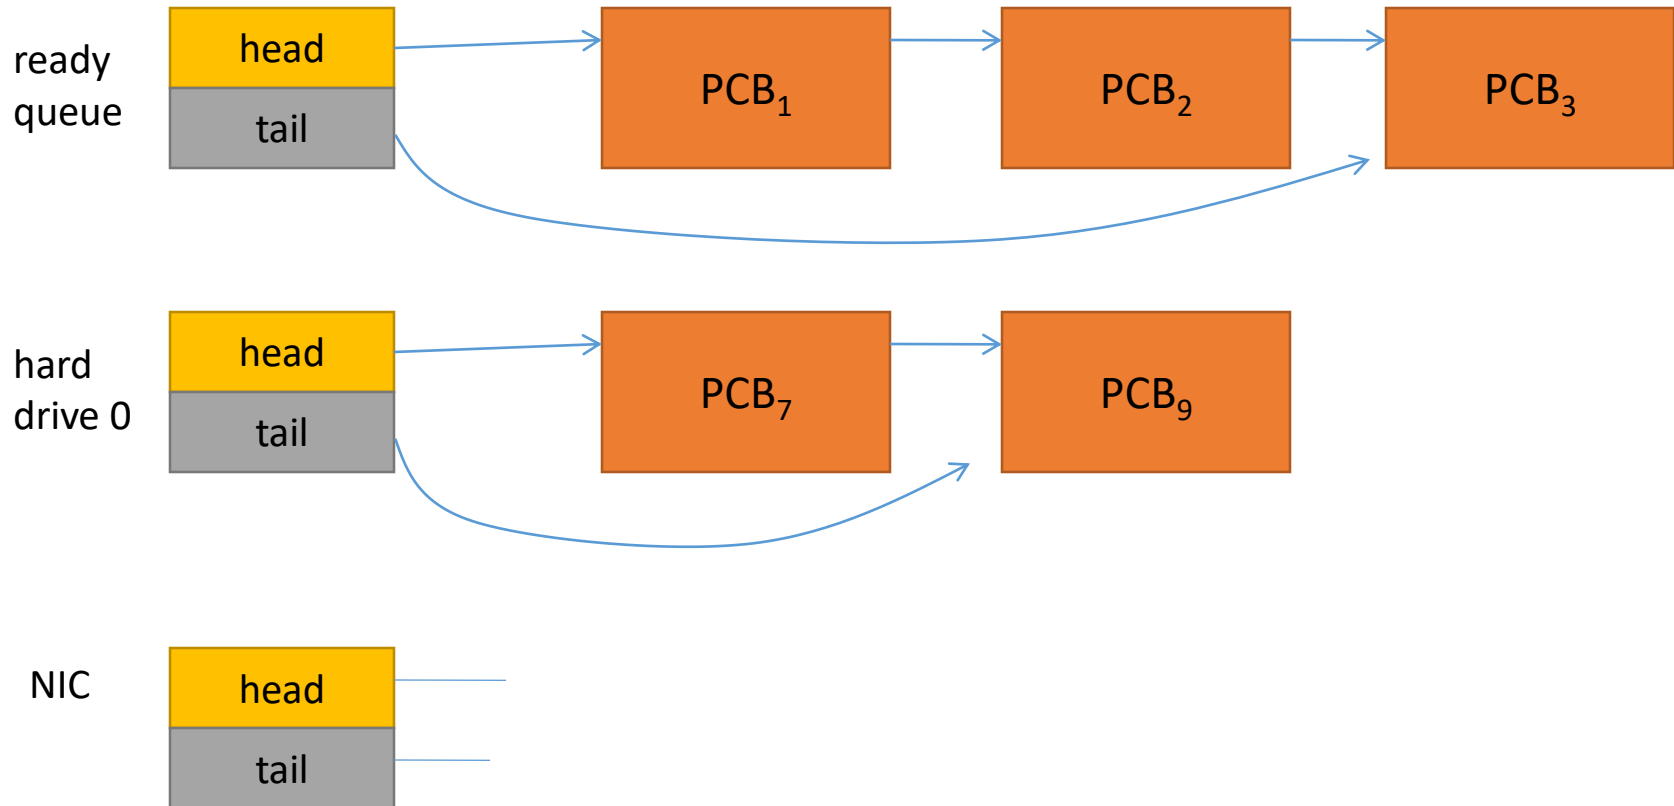
- Register values stored in  $PCB_i$
- PC stored in  $PCB_i$
- $PCB_i$  state changed to *ready*
- Kernel schedules next process
  - May be same process is only one is currently executing
  - Still follows same rules

When process  $P_j$  returned to CPU,

- Register values restored from  $PCB_j$
- $PCB_j$  state changed to *running*
- CPU mode changed to User-mode
- PC copied from  $PCB_j$  to CPU
  - Jumps back to process
  - Last thing to happen

- OS maintains multiple queues to manage running processes.
  - each lists PCBs of processes
  - Job queue
  - Ready queue
  - Device queues
- OS moves processes to different queues depending on their status.

# Process Queues in Memory



- During boot, one process is created
  - PID is 1
  - Usually named “init”
- All other processes are *descendents* of init
- Processes form hierarchy:
  - parents & children
  - all children have one parent
  - parents can have any number of children

- Address space:
  - parent and child process may share address space
  - create a new address space for child




- System call to create new process:

`fork()`

- takes no parameters
- creates copy of current proc
- new PCB, new memory, same content
- memory copy: old to new
- PCB copy: old to new
- new PID

# Process Tree in UNIX



Stack	i	?
	j	?
Data	g	?
Code	LOAD R2, 100	
	STOR R2, j	
	LOAD R3, 5	
	STOR R3, g	
	CALL fork	
	STOR RV, i	
PCB	PID	<b>296</b>
	R1	?
	R2	?
	R3	?
	PC	<b>0</b>

```
int g;
```

```
int main() {
```

```
    int i, j;
```

```
    j = 100;
```

```
    g = 5;
```

```
    i = fork();
```

```
    printf("%d:%d:%d\n", g, i, j);
```

```
    return 0;
```

```
}
```

Stack	i	?
	j	?
Data	g	?
Code	LOAD R2, 100	
	STOR R2, j	
	LOAD R3, 5	
	STOR R3, g	
	CALL fork	
	STOR RV, i	
PCB	PID	<b>296</b>
	R1	?
	R2	<b>100</b>
	R3	?
	PC	<b>1</b>

```
int g;
```

```
int main() {
```

```
    int i, j;
```

```
    j = 100;
```

```
    g = 5;
```

```
    i = fork();
```

```
    printf("%d:%d:%d\n", g, i, j);
```

```
    return 0;
```

```
}
```

Stack	i	?
	j	<b>100</b>
Data	g	?
Code	LOAD R2, 100	
	<b>STOR R2, j</b>	
	LOAD R3, 5	
	STOR R3, g	
	CALL fork	
	STOR RV, i	
PCB	PID	<b>296</b>
	R1	?
	R2	<b>100</b>
	R3	?
	PC	<b>2</b>

```
int g;
```

```
int main() {
```

```
    int i, j;
```

```
    j = 100;
```

```
    g = 5;
```

```
    i = fork();
```

```
    printf("%d:%d:%d\n", g, i, j);
```

```
    return 0;
```

```
}
```

Stack	i	?
	j	<b>100</b>
Data	g	?
Code	LOAD R2, 100	
	STOR R2, j	
	LOAD R3, 5	
	STOR R3, g	
	CALL fork	
	STOR RV, i	
PCB	PID	<b>296</b>
	R1	?
	R2	<b>100</b>
	R3	<b>5</b>
	PC	<b>3</b>

```
int g;
```

```
int main() {
```

```
    int i, j;
```

```
    j = 100;
```

```
    g = 5;
```

```
    i = fork();
```

```
    printf("%d:%d:%d\n", g, i, j);
```

```
    return 0;
```

```
}
```

Stack	i	?
	j	<b>100</b>
Data	g	<b>5</b>
Code	LOAD R2, 100	
	STOR R2, j	
	LOAD R3, 5	
	<b>STOR R3, g</b>	
	CALL fork	
	STOR RV, i	
PCB	PID	<b>296</b>
	R1	?
	R2	<b>100</b>
	R3	<b>5</b>
	PC	<b>4</b>

```
int g;
```

```
int main() {
```

```
    int i, j;
```

```
    j = 100;
```

```
    g = 5;
```

```
    i = fork();
```

```
    printf("%d:%d:%d\n", g, i, j);
```

```
    return 0;
```

```
}
```

Stack	i	?
	j	<b>100</b>
Data	g	<b>5</b>
Code	LOAD R2, 100	
	STOR R2, j	
	LOAD R3, 5	
	STOR R3, g	
	CALL fork	
	STOR RV, i	
PCB	PID	<b>296</b>
	R1	?
	R2	<b>100</b>
	R3	<b>5</b>
	PC	<b>5</b>

```

int g;

int main() {
    int i, j;
    j = 100;
    g = 5;
    i = fork();
    printf("%d:%d:%d\n", g, i, j);
    return 0;
}

```



# Process P<sub>296</sub> and Process P<sub>321</sub>

Stack	i	?
	j	<b>100</b>
Data	g	<b>5</b>
Code	LOAD R2, 100	
	STOR R2, j	
	LOAD R3, 5	
	STOR R3, g	
	CALL fork	
	STOR RV, i	
PCB	PID	<b>296</b>
	R1	?
	R2	<b>100</b>
	R3	<b>5</b>
	PC	<b>5</b>

← parent  
child →

1 call, 2  
returns

Independent  
processes

Execute in  
any order

Can be  
interleaved

Assume child  
executes 1<sup>st</sup>

Stack	i	?
	j	<b>100</b>
Data	g	<b>5</b>
Code	LOAD R2, 100	
	STOR R2, j	
	LOAD R3, 5	
	STOR R3, g	
	CALL fork	
	STOR RV, i	
PCB	PID	<b>321</b>
	R1	?
	R2	<b>100</b>
	R3	<b>5</b>
	PC	<b>5</b>

- Local variables survive
- Global variables survive
- File descriptor table survives
- PCB
  - Registers survive
  - Stack pointer doesn't survive
  - PC doesn't survive
  - PID doesn't survive

# Process P<sub>296</sub> and Process P<sub>321</sub>

Stack	i	?
	j	<b>100</b>
Data	g	<b>5</b>
Code	LOAD R2, 100	
	STOR R2, j	
	LOAD R3, 5	
	STOR R3, g	
	CALL fork	
	STOR RV, i	
PCB	PID	<b>296</b>
	R1	?
	R2	<b>100</b>
	R3	<b>5</b>
	PC	<b>5</b>

Child didn't  
call fork()

Child gets a  
zero return  
value

0 => child

Quantum  
Expires

Parent's turn

Stack	i	<b>0</b>
	j	<b>100</b>
Data	g	<b>5</b>
Code	LOAD R2, 100	
	STOR R2, j	
	LOAD R3, 5	
	STOR R3, g	
	CALL fork	
	STOR RV, i	
PCB	PID	<b>321</b>
	R1	?
	R2	<b>100</b>
	R3	<b>5</b>
	PC	<b>6</b>

# Process P<sub>296</sub> and Process P<sub>321</sub>

Stack	i	<b>321</b>
	j	<b>100</b>
Data	g	<b>5</b>
Code	LOAD R2, 100	
	STOR R2, j	
	LOAD R3, 5	
	STOR R3, g	
	CALL fork	
	STOR RV, i	
PCB	PID	<b>296</b>
	R1	<b>?</b>
	R2	<b>100</b>
	R3	<b>5</b>
	PC	<b>6</b>

Parent did  
call fork()

Parent gets  
PID of child

!0 => parent

Quantum  
Expires

Child's turn

Stack	i	<b>0</b>
	j	<b>100</b>
Data	g	<b>5</b>
Code	LOAD R2, 100	
	STOR R2, j	
	LOAD R3, 5	
	STOR R3, g	
	CALL fork	
	STOR RV, i	
PCB	PID	<b>321</b>
	R1	<b>?</b>
	R2	<b>100</b>
	R3	<b>5</b>
	PC	<b>6</b>

# Process P<sub>296</sub> and Process P<sub>321</sub>

Stack	i	<b>321</b>
	j	<b>100</b>
Data	g	<b>5</b>
Code	LOAD R2, 100	
	STOR R2, j	
	LOAD R3, 5	
	STOR R3, g	
	CALL fork	
	STOR RV, i	
PCB	PID	<b>296</b>
	R1	<b>?</b>
	R2	<b>100</b>
	R3	<b>5</b>
	PC	<b>?</b>

## OUTPUT

5:0:100

5:321:100

## OR

5:321:100

5:0:100

Stack	i	<b>0</b>
	j	<b>100</b>
Data	g	<b>5</b>
Code	LOAD R2, 100	
	STOR R2, j	
	LOAD R3, 5	
	STOR R3, g	
	CALL fork	
	STOR RV, i	
PCB	PID	<b>321</b>
	R1	<b>?</b>
	R2	<b>100</b>
	R3	<b>5</b>
	PC	<b>?</b>

- Program has 1 `printf()` call
- Output has 2 lines
- 1 process (parent) before `fork()`
- 2 processes (parent, child) after `fork()`
- Parent gets PID of child returned from `fork()`
- Child gets 0 returned from `fork()`
- 1 call to `fork()` , 2 returns
- Now we can clone process

```
int main(int argc, char ** argv)
{
    pid_t pid;
    printf("Output line 1\n");
    pid = fork();
    printf("Output line 2 with pid = %d\n", pid);
    return 0;
}
```

What is the output?

Assume parent is PID 94 and child is PID 223.

```
int main(int argc, char ** argv)
{
    pid_t pid;
    printf("Output line 1\n");
    pid = fork();
    printf("Output line 2 with pid = %d\n", pid);
    return 0;
}
```


What is the output?

Assume parent is PID 94 and child is PID 223.

Output line 1

Output line 2 with pid = 223

Output line 2 with pid = 0



Parent printed  
first



```
int main(int argc, char ** argv)
{
    pid_t pid;
    printf("Output line 1\n");
    pid = fork();
    printf("Output line 2 with pid = %d\n", pid);
    return 0;
}
```


What is the output?

Assume parent is PID 94 and child is PID 223.

Output line 1

Output line 2 with pid = 0

Output line 2 with pid = 223



Child printed  
first

```
int main(int argc, char ** argv)
{
    pid_t pid;
    doCommonEntryCode();
    pid = fork();
    if (pid) {
        doParentCode();
    } else {
        doChildCode();
    }
    doDuplicatedExitCode();
    return 0;
}
```

- We can clone process
- We can create clone of `myshell`
- We want to create `ls` instead
- How?
- Try the `exec()` family of system calls

```
int execl(const char *path, const char *arg, ...);  
int execvp(const char *file, const char *arg, ...);  
int execle(const char *path, const char *arg ,  
    ..., char * const envp[]);  
int execv(const char *path, char *const argv[]);  
int execvp(const char *file, char *const argv[]);
```

The `exec` family of functions replaces the current process image with a new process image.

The initial argument for these functions is the pathname of a file which is to be executed.

The `const char *arg` and subsequent ellipses in the `execl()`, `execvp()`, and `execle()` functions can be thought of as `arg0, arg1, ..., argn`. They describe a list of one or more pointers to null-terminated strings that represent the argument list available to the executed program. The first argument should point to the file name associated with the file being executed. The list of arguments must be terminated by a `NULL` pointer, and, since these are variadic functions, this pointer must be cast `(char *) NULL`.

```
#include <stdio.h>

int main(int argc, char ** argv) {
    int i;
    for (i = 0; i < argc; i++)
        printf("argv[%d] = |%s|\n",
               i, argv[i]);
    return 0;
}
```

```
gcc -g -Wall printargs.c -o pa
```

```
> pa
```

```
argv[0] = |pa|
```

```
> pa -m file
```

```
argv[0] = |pa|
```

```
argv[1] = |-m|
```

```
argv[2] = |file|
```

```
>
```



```
pid = fork();  
if (pid) {  
    printf("parent: pid = %d\n", pid);  
} else {  
    execl("./pa", "pa", NULL);  
    printf("child: pid = %d\n", pid);  
    exit(1);  
}  
printf("more output\n");
```

# Process P<sub>296</sub> and Process P<sub>321</sub>

Stack	i	?
	j	<b>100</b>
Data	g	<b>5</b>
Code	LOAD R2, 100	
	STOR R2, j	
	LOAD R3, 5	
	STOR R3, g	
	CALL fork	
	STOR RV, i	
PCB	PID	<b>296</b>
	R1	?
	R2	<b>100</b>
	R3	<b>5</b>
	PC	<b>5</b>

after fork

2 copies of  
parent

Stack	i	?
	j	<b>100</b>
Data	g	<b>5</b>
Code	LOAD R2, 100	
	STOR R2, j	
	LOAD R3, 5	
	STOR R3, g	
	CALL fork	
	STOR RV, i	
PCB	PID	<b>321</b>
	R1	?
	R2	<b>100</b>
	R3	<b>5</b>
	PC	<b>5</b>

# Process P<sub>296</sub> and Process P<sub>321</sub>

Stack	i	?
	j	<b>100</b>
Data	g	<b>5</b>
Code	LOAD R2, 100	
	STOR R2, j	
	LOAD R3, 5	
	STOR R3, g	
	CALL fork	
	STOR RV, i	
PCB	PID	<b>296</b>
	R1	?
	R2	<b>100</b>
	R3	<b>5</b>
	PC	<b>5</b>

after execl()

← 1 parent  
executing  
original code

1 child  
executing pa →

Stack		
Data		
Code	exec for pa	
	...	
	...	
	...	
	...	
	...	
PCB	PID	<b>321</b>
	R1	?
	R2	?
	R3	?
	PC	<b>0</b>

# Process P<sub>296</sub> and Process P<sub>321</sub>

Stack	i	?
	j	100
Data	g	5
Code	LOAD R2, 100	
	STOR R2, j	
	LOAD R3, 5	
	STOR R3, g	
	CALL fork	
	STOR RV, i	
PCB	PID	296
	R1	?
	R2	100
	R3	5
	PC	5

after execl()

Stack, data,  
initialized

Code for pa  
overwrites  
code for  
parent

PCB  
initialized

Stack		
Data		
Code	exec for pa	
	...	
	...	
	...	
	...	
	...	
PCB	PID	321
	R1	?
	R2	?
	R3	?
	PC	0

```
parent: pid = 1093  
argv[0] = |pa|  
more output
```

**OR**

```
parent: pid = 1093  
more output  
argv[0] = |pa|
```

output determined by  
quantum expiration  
times

```
pid = fork();  
if (pid) {  
    printf("parent: pid = %d\n", pid);  
} else {  
    execl("./pa", "pa", "-m", "file",  
          NULL);  
    printf("child: pid = %d\n", pid);  
    exit(1);  
}  
printf("more output\n");
```

```
parent pid = 1093  
argv[0] = |pa|  
more output  
argv[1] = |-m|  
argv[2] = |file|
```

output still  
determined by  
quantum expiration  
times

# Why so many `exec()`s?

Allow argument passing in different forms

- `execl()` – give path to exe, separate args
- `execvp()` – give name of exe, separate args
- `execv()` – give path to exe, arg vector
- `execvp()` – give name of exe, arg vector

Project 1: Check the values produced from the first part of the project. What form are they in? Is that useful?



# What happens after `exec()` ?

- Local variables don't survive
- Global variables don't survive
- File descriptor table survives
- PCB
  - Registers don't survive
  - Stack pointer doesn't survive
  - PC doesn't survive
  - PID survives

- A process is a running program in memory.
- The OS creates a PCB for each process.
- The process goes through different stages during its life cycle.
  - may be performing IO, then it is blocked
  - may be running or waiting or terminating or starting
- A process is created through `fork()`.
  - except for the first process, called `init`
- System call `exec` loads the code of a process into memory.