



UNIVERSITY *of* WEST FLORIDA

COP4634: Systems & Networks I

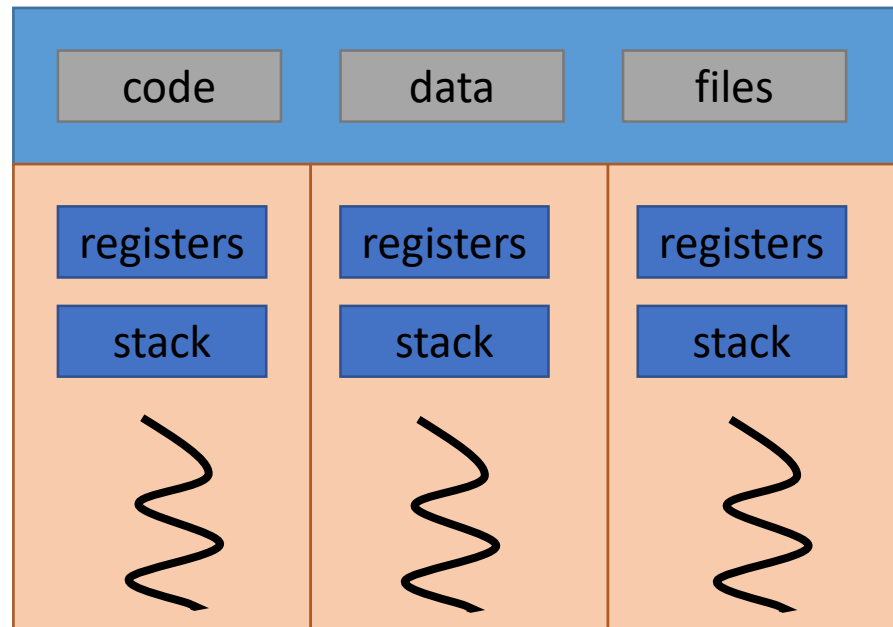
Threads

- Definition: A thread (= job) is a lightweight process, representing a single unit of program execution within a process.
 - shares code, data & file descriptor table
 - has individual stack, registers & thread control block (TCB)
- A process may run multiple threads concurrently.

- Two types of threads
 - User-level threads
 - Kernel threads
- Thread creation similar to process creation except:
 - Processes begins execution at main().
 - Thread begins execution at specified function.
 - Creation call identifies initial function.

Multithreaded Programming

- Multiple threads may be executed simultaneously in a single process.
- Threads share code, global data, & open files of the process.



Process – Threads Differences


- Threads share process resources, processes share computer resources.
 - ▶ OS maintains per thread
 - ▶ program counter,
 - ▶ stack
 - ▶ state
 - ▶ registers
 - ▶ OS maintains per processes
 - ▶ an address space,
 - ▶ list of open files,
 - ▶ child process,
 - ▶ signals and signal handlers
 - ▶ accounting

- Threads are more responsive than processes.
- Threads share resources.
- Threads are more economical than processes.

- Client – Server program
 - server handles each client request in a separate thread
 - server constructs a pool of threads to distribute work
- Web Browser
 - multiple browser windows downloading Web pages simultaneously
 - may not be multi-threaded but multi-processed


Multithreading Models

- Many-to-one model: many user-level threads maps to one kernel thread.
 - if a single thread makes a blocking call, all threads will be blocked




Multithreading Models (cont.)

- One-to-one model: a user-level thread maps to a single kernel thread.
 - provides more concurrency
 - more overhead to create a new user level thread
 - can only create as many user level threads as there are kernel threads available



Multithreading Models (cont.)

- Many-to-many model: many user-level threads map to many kernel threads.
 - combines advantages of one-to-one and many-to-many model



Example – Web Server

- Parent waits for http request from network
- Parent creates child thread to handle request
- Child thread:
 - reads disk to find http page
 - sends http page to client
 - exits
- Parent thread:
 - waits for next request

Review: Process P₂₉₆ Before fork()

Stack	i	?
	j	100
Data	g	5
Code	LOAD R2, 100	
	STOR R2, j	
	LOAD R3, 5	
	STOR R3, g	
	CALL fork	
	STOR RV, i	
PCB	PID	296
	R1	?
	R2	100
	R3	5
	PC	5

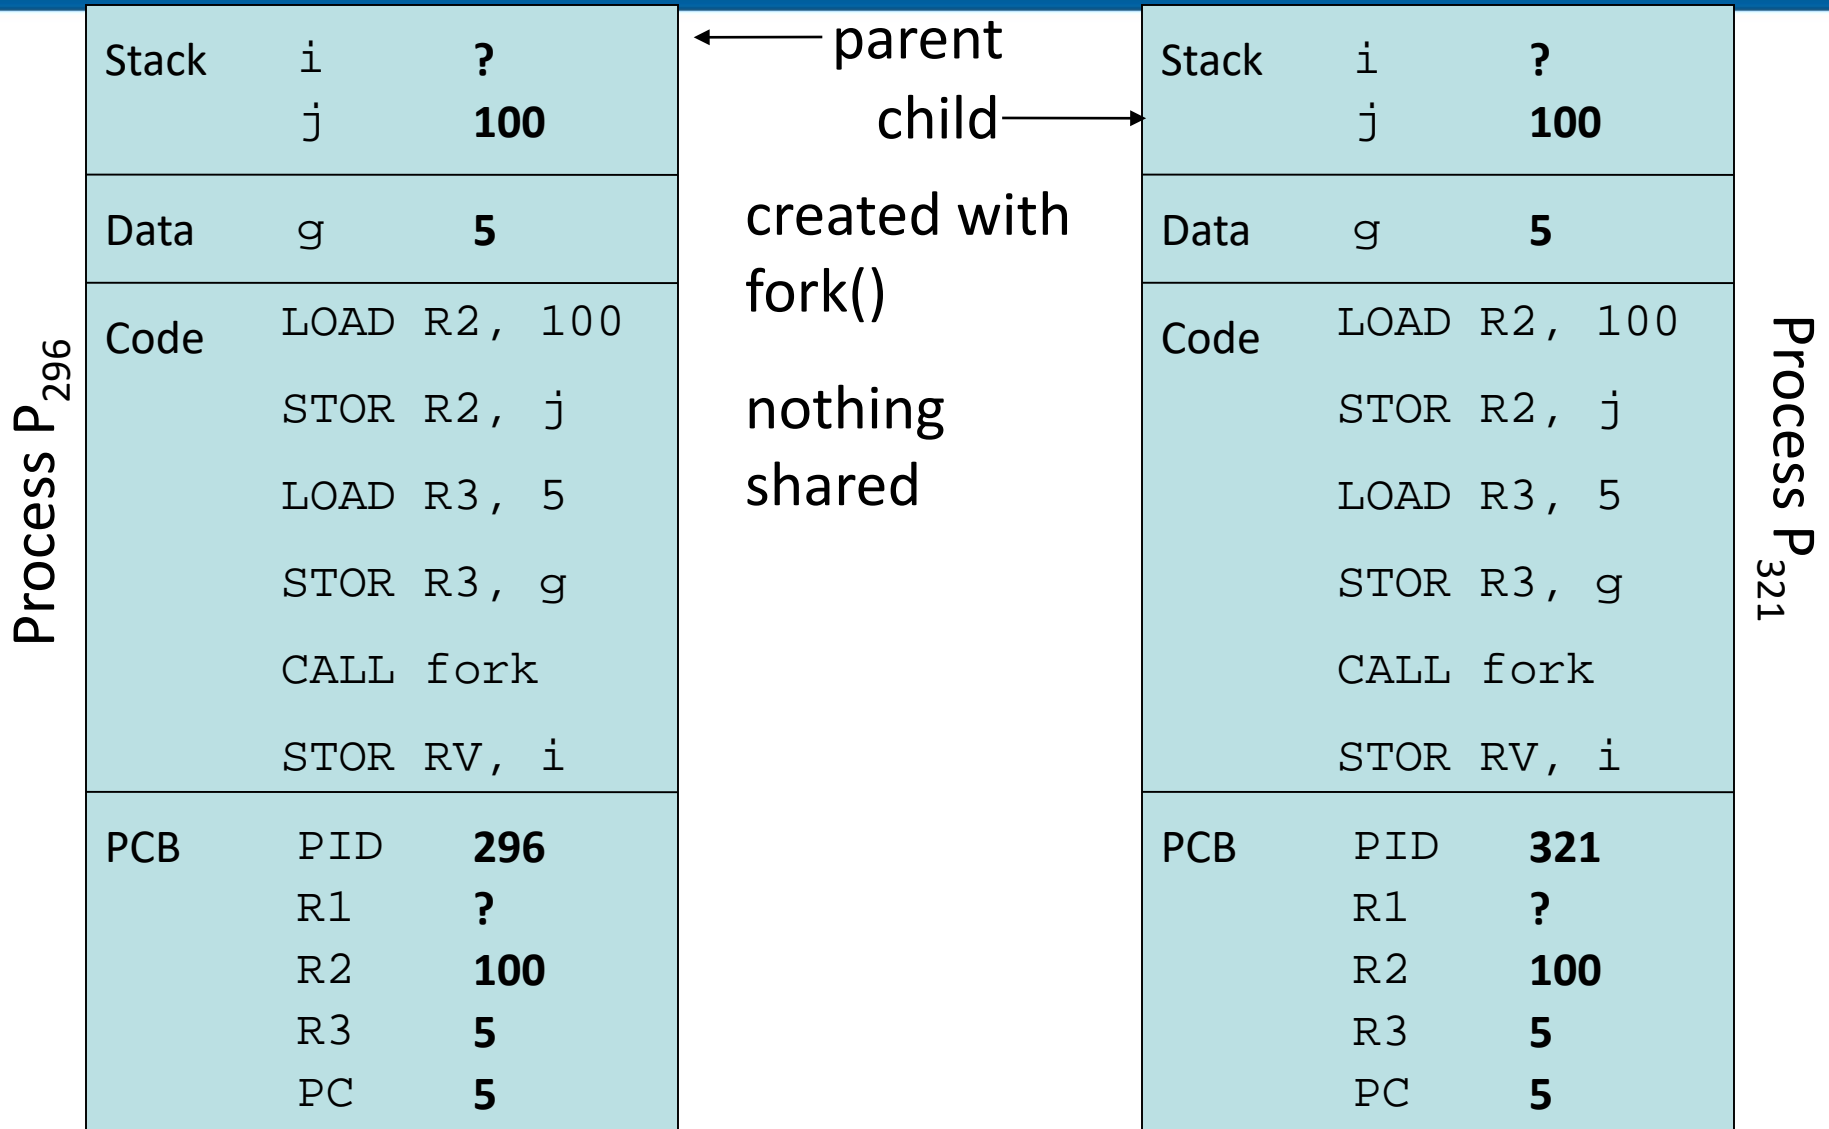
```

int g;

int main() {
    int i, j;
    j = 100;
    g = 5;
    i = fork();
    printf("%d:%d:%d\n", g, i, j);
    return 0;
}

```

Review: Processes P₂₉₆ After fork()



An indication of

- where is current execution state
 - Register values in PCB per thread (aka TCB)
- how this point was reached
 - Runtime stack per thread
- what should be done next
 - PC per thread

Process – Another View of Process P₂₉₆

Thread
(of control)

“single-
threaded
process”

OR

“single-
threaded task”

OR

process

Data	g	5	Code	LOAD R2, 100
File Descriptor Table	stdin			STOR R2, j
	stdout			LOAD R3, 5
	stderr			STOR R3, g
				CALL fork
				STOR RV, i
Stack	i	?		
	j	100		
TCB	TID	296		
	R1	?		
	R2	100		
	R3	5		
	PC	5		

“multi-threaded task”

11

- Define a function for the thread
 - AR for function pushed onto thread stack
 - Thread terminates when function returns
 - must be return type `void *`
 - must accept one param of type `void *`

```
void * threadMain(void *threadParam);  
void * doCompute(void *records);  
void * sortThis(void *array);
```

What is `void *`?

*type **

pointer to type

`void`

anything - or nothing

`void *`

pointer to anything

can be assigned any pointer

```
void * vp;
```

```
char * cp;
```

```
int * ip;
```

```
double * dp;
```

```
...
```

```
vp = cp;
```

```
vp = ip;
```

```
vp = dp;
```

What can be param?

What about other types? Send pointer.

Must be typecasted to `void *`

Compiler will complain otherwise

```
double d;
```

```
int i;
```

```
...
```

```
threadFunctionA( (void *)&d );
```

```
threadFunctionB( (void *)&i );
```

Typecasting doesn't change data;
it only shuts up compiler

```
long val1, val2;  
char *cp;  
val1 = 123456789;  
cp = (char *)val1;  
val2 = (long)cp;  
printf("%ld\t%ld\n", val1, val2);
```

1. Typecast what you are passing

```
int i = 123456;
```

```
threadFunctionB( (void *)&i );
```



do it

2. Typecast what you received

```
void* threadFunctionB(void *param)  
{
```

```
    int x = *(int *)param;
```

```
    printf("%d\n", x);
```



undo it

Example Problem in C

```
void * threadFunctionB( void * p) {  
    int n = *(int *)p;  
    printf( "%d\n", n);  
}
```

```
int i;    // a global variable  
i = 1;    // initialize variable  
create_thread(threadFunctionB((void *)&i ));  
i = 2;  
create_thread(threadFunctionB((void *)&i ));  
i = 3;
```

One Thread

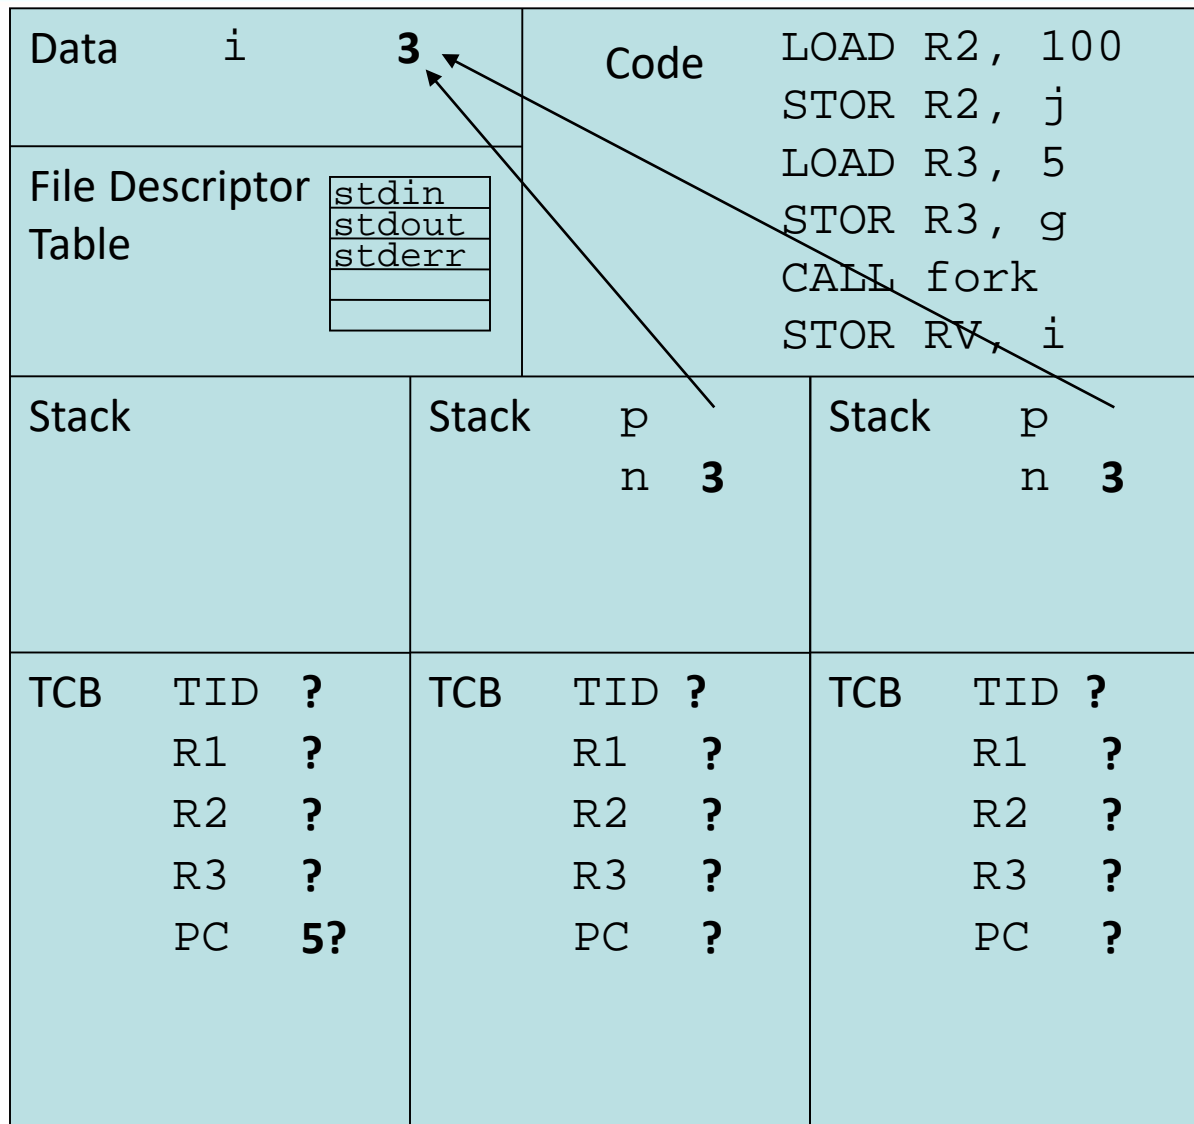
Data			i	1	Code	LOAD R2, 100					
						STOR R2, j					
						LOAD R3, 5					
						STOR R3, g					
						CALL fork					
File Descriptor Table			<table><tr><td>stdin</td></tr><tr><td>stdout</td></tr><tr><td>stderr</td></tr><tr><td></td></tr><tr><td></td></tr></table>		stdin	stdout	stderr				STOR RV, i
stdin											
stdout											
stderr											
Stack											
TCB					TID	?					
					R1	?					
					R2	?					
					R3	?					
					PC	5?					

Two Threads

Data i 2			Code	LOAD R2, 100 STOR R2, j LOAD R3, 5 STOR R3, g CALL fork STOR RV, i					
File Descriptor Table									
<table><tr><td>stdin</td></tr><tr><td>stdout</td></tr><tr><td>stderr</td></tr><tr><td> </td></tr><tr><td> </td></tr></table>			stdin	stdout	stderr				
stdin									
stdout									
stderr									
Stack		Stack p ? n ?							
TCB	TID	?	TCB TID ?						
	R1	?			R1 ?				
	R2	?		R2 ?					
	R3	?		R3 ?					
	PC	5?		PC ?					

31

Shared Memory Location



- IEEE 1003.1c POSIX API standard.
- Requirement for interface.
- Not a requirement for implementation.

```
pthread_create( ... );
```

```
pthread_exit( RV );
```

```
pthread_join( threadID );
```

- Create a `void *function(void *)`
- Create a thread / call the function
 - Typecast the argument (`void *`)
- Function (child)
 - undo the argument typecasting
 - perform child thread requirements
 - return from the child thread
- Main (parent)
 - perform main thread requirements
 - wait for child thread to return

Another Example (in C)

```
int g;  /* global variable */

void * tFunc(void *param){
    int i;
    char *str = (char *)param;
    for( i=0; i<3; i++ ){
        printf("%s [%d:%d]\n",str, i, g);
        g++;
    }
    pthread_exit(0);
}
```

Calling threadfunc()

```
int main(int argc, char ** argv){
    pthread_t tidA, tidB;

    g = 1;
    printf("prethreads [%d]\n", g);

    pthread_create(&tidA, NULL, tFunc, (void *)"threadA");
    pthread_create(&tidB, NULL, tFunc, (void *)"threadB");

    printf("postthreads [%d]\n", g);

    pthread_join(tidA, NULL);
    pthread_join(tidB, NULL);

    printf("main done [%d]\n", g);

    return 0;
}
```

- Passing thread function similar to C but simpler.

```
int main(int argc, char ** argv){
    g = 1;
    printf("prethreads [%d]\n", g);

    // create and start threads
    std::thread first (tFunc, "threadA");
    std::thread second (tFunc, "threadB");

    printf("postthreads [%d]\n", g);

    // join threads
    first.join();
    second.join();

    printf("main done [%d]\n", g);

    return 0;
}
```

One Possible Execution

```
prethreads [1]
postthreads [1]
threadA [0:1]
threadA [1:2]
threadB [0:3]
threadB [1:4]
threadB [2:5]
threadA [2:6]
main done [7]
```

Assumptions

- Main thread runs before any child
- Main thread waits for new threads to terminate
- Thread A prints 2 lines before quantum expires
- Thread B runs to completion
- Thread A completes


- Library or Kernel
- Library – original method
 - thread is executed in user mode
 - kernel unaware of threads
 - processes are scheduled, not threads
 - fast to create/switch/destroy
 - `yield()` gives CPU to next process
 - quantum shared among threads executed within process

- Assumptions
 - 10 tick context switch in kernel
 - 2 tick context switch in library
 - 100 tick quantum
- We already have 1 process in system (D)
- New process (P) assumptions
 - arrives at time 33
 - composed of 1 thread
 - needs 500 ticks to complete computation
 - quantum shared equally among threads

Time – 1 Library Thread

•	0	D							
•	100	K		440	D		880	D	
•	110	P 100		540	K		980	K	
•	210	K		550	P 300		990	P 500	
•	220	D		650	K		1090	K	
•	320	K		660	D		1100	D	
•	330	P 200		760	K		1200	K	
•	430	K		770	P 400		1210	D	
				870	K		1310	K	

Time – 1 Library Thread



Timeline of what process is on the CPU

Assumes P arrives at time 33 (*new*)

P goes to end of ready list (*ready*)

D already on CPU

- t_a arrival time
- t_e execution start time
- t_d departure time
- Response = $t_e - t_a$
 - fastest possible user-observable reaction
- Turnaround = $t_d - t_a$
 - press enter until prompt returns
 - enter system until exit system


Stats from Examples

	T_a	T_e	T_d	Resp	T/A
1 thread	33	110	1090	77	1057
5 library threads					
5 kernel threads					

Time – 5 Library Threads

0	D	440	D	880	D
100	K	540	K	980	K
110	P	550	P	990	P
210	K	650	K	1090	K
220	D	660	D	1100	D
320	K	760	K	1200	K
330	P	770	P	1210	P
430	K	870	K	1264	K

Time – 5 Library Threads




Time – 5 Library Threads

0	D		440	D		880	D
100	K		540	K		980	K
110	P		550	P		990	P
210	K	90	650	K	270	1090	K
220	D		660	D		1100	D
320	K		760	K		1200	K
330	P		770	P		1210	P
430	K	180	870	K	360	1264	K
							500

Ready List: D, P

Time – 5 Library Threads



Stops at 500

Continues 18/2

Doesn't use all of
quantum

Stats from Examples

	T_a	T_e	T_d	Resp	T/A
1 thread	33	110	1090	77	1057
5 library threads	33	110	1264	77	1231
5 kernel threads					

- Library or Kernel
- Kernel – newer method
 - implemented in kernel code
 - kernel aware of threads (TID & PID)
 - threads are scheduled
 - slow to create/switch/destroy
 - `yield()` gives CPU to next thread
 - quantum explicitly for thread

- Assumptions
 - 10 tick context switch in kernel
 - 100 tick quantum
- We already have 1 process in system (D)
- New process (P) assumptions
 - composed of 1 thread
 - needs 500 ticks to complete computation

Time – 1 Kernel Thread

0	D		440	D		880	D
100	K		540	K		980	K
110	P	100	550	P	300	990	P
210	K		650	K		1090	K
220	D		660	D		1100	D
320	K		760	K		1200	K
330	P	200	770	P	400	1210	D
430	K		870	K		1310	K

Stats from Examples

	T_a	T_e	T_d	Resp	T/A
1 thread Same	33	110	1090	77	1057
5 library threads	33	110	1264	77	1234
5 kernel threads					

Time – 5 Kernel Threads

0	D	440	P_3	400	880	D
100	K	540	K		980	K
110	P_0	550	P_4	500	990	D
210	K	650	K		1090	K
220	P_1	660	D		1100	D
320	K	760	K		1200	K
330	P_2	770	D		1210	D
430	K	870	K		1264	K

Ready List: D, P_0 , P_1 , P_2 , P_3 , P_4

Stats from Examples

	T_a	T_e	T_d	Resp	T/A
1 thread	33	110	1090	77	1057
5 library threads	33	110	1264	77	1231
5 kernel threads	33	110	650	77	617

- Library threads
 - kernel schedules process
 - thread initiates I/O, process blocks
 - no other thread can run
- Kernel threads
 - kernel schedules threads
 - thread initiates I/O, thread blocks
 - doesn't effect other threads

- Which is better for web server implementation?
 - library threads?
 - kernel threads?
 - multiple processes (1 thread each)?
- Why?

- Threads are units of executions within a process.
- Multiple threads may be executed simultaneously within a process.
- Threads have the states ready, running, and blocked.
- Threads can be kernel or user-level threads.
- Kernel threads: OS executes threads, not processes.
- Threads may need to be synchronized to avoid race conditions.