



UNIVERSITY *of* WEST FLORIDA

COP4634: Systems & Networks I

Synchronization

- Concurrent access to shared data may result in data becoming inconsistency.
- Maintaining data consistency while supporting concurrent processing requires additional OS mechanisms.

Example: Concurrent Data Access

```
int g;
void *tFunc(void *param){
    char *str = (char *)param;
    int loc = 0;
    g += 10;
    loc = g + 5;
    printf("%s: g %d: loc %d\n", str, g, loc);
    pthread_exit(0);
}
int main(int argc, char **argv){
    pthread_t tid1, tid2;
    g = 10;
    pthread_create(&tid1, NULL, tFunc, (void *)"ONE");
    pthread_create(&tid2, NULL, tFunc, (void *)"TWO");
    pthread_join(tid1, NULL);
    pthread_join(tid2, NULL);
    return 0;
}
```

```
int g;
void *tFunc(void *param){
    char *str = (char *)param;
    int loc = 0;
    g += 10;
        LOAD R3, 10
        LOAD R4, g
        ADD R5,R4,R3
        STOR R5, g
    loc = g + 5;
        LOAD R3, 5
        LOAD R4, g
        ADD R5,R4,R3
        STOR R5, loc
    printf("%s: g %d: loc %d\n", str, g, loc);
        LOAD R3, str
        LOAD R4, g
        LOAD R5, loc
    ...
    pthread_exit(0);
}
```

g

ONE TWO

R3

R4

R5

loc

Race Condition:

the final outcome of a result is determined by the order of process execution

Critical Section: a code section in a process that accesses shared data

```
void *tFunc(void *param){
    char *str = (char *)param;
    int loc = 0;
    g += 10;
    loc = g + 5;
    printf("%s: g %d: loc %d\n", str, g, loc);
    pthread_exit(0);
}
```

- *Mutual Exclusion:*
 - only one thread at a time can be in their critical section
- *Progress:*
 - if no thread is in their critical section, other processes should be allowed to enter
- *Bounded waiting:*
 - there must be a bound on the time a thread waits to enter its critical section

- Threads must proceed through an entry and exit sections.
- Threads must not die or quit inside a critical section.
- Threads may be context switched inside a critical section.
 - a context switch is different from an exit because

Bad Solution 1

```
int turn;  
  
1 while (1) {  
2     ...  
3     while (turn != i);  
4     CS();  
5     turn = j;  
6     ...  
7 }
```

- CS – Critical Section
- Mutual exclusion (mutex)
- i – pid of currently executing process
- j – pid of “other” process


```
1 while(1) {
2     ...
3     while (turn != i);
4     CS();
5     turn = j;
6     ...
7 }
```

```
turn    P0    P1
```

wantCS	0	1
--------	---	---

wantCS	0	1
--------	---	---

```
while (1) {
```

```
    ...
```

```
    ENTRY_CODE ();
```

```
    CS ();
```

```
    EXIT_CODE ();
```

```
    ...
```

```
}
```

Controls access to CS

Insures MUTEX

Allows next P to enter CS

Assume no crashes for
now

- Once started, must complete
- Can't be interrupted in the middle
- Completes or doesn't start
- Statements (generally) NOT atomic
- Instructions are atomic

- Peterson's and Bakery Algorithm
 - work but are difficult to generalize to all sorts of computing problems requiring synchronization
- Synchronization Hardware
 - atomic instructions that test and set a boolean value at the same time

- User-level implementation:
 - bugs can be easily introduced
 - synchronization code can be (un)intentionally skipped
 - requires busy wait loop
- What we want:
 - kernel implementation
 - user can use
 - user cannot alter
 - efficient implementation

- Kernel implementation of MUTEX
- Has two functions
 - Acquire() requests lock, only returns when lock is given to process
 - Release() gives lock back to kernel to reallocate to another process
- Presented as “Class/Object” for clarity
- Really implemented in C

```
class LOCK {  
    int status = FREE;  
}  
LOCK::Acquire() {  
    Disable_Interrupts();  
    while (status == BUSY) {  
        Enable_Interrupts();  
        Disable_Interrupts();  
    }  
    status = BUSY;  
    Enable_Interrupts();  
}
```



```
class LOCK {  
    int status = FREE;  
}  
LOCK::Acquire() {  
    ...  
}  
LOCK::Release() {  
    Disable_Interrupts();  
    status = FREE;  
    Enable_Interrupts();  
}
```

Why not Perfect Solution?

- Busy wait loop still present
 - Eats CPU time while waiting
- High priority thread waiting?
 - Holder (low priority) never runs
- Waiting thread could be “infinitely unlucky”
 - How could this be?

```
class LOCK {  
    int status = FREE;  
}  
LOCK::Acquire() {  
    Disable_Interrupts();  
    while (status == BUSY) {  
        Enable_Interrupts();  
        Disable_Interrupts();  
    }  
    status = BUSY;  
    Enable_Interrupts();  
}
```

P0
interrupted
here



P1 Release(),
Acquire(),
status is FREE

- Add “Waiting List for Lock” (queue)
- Need `enqueue()` and `dequeue()` for list
- Must be able to put “self” on list

Idea:

If lock is BUSY,
 put self on waiting list
 go to sleep

When releasing lock,
 wake-up a sleeping thread

```
class LOCK{
    int status = FREE;
    queue waiting;
}
LOCK::Acquire() {
    Disable_Interrupts();
    while (status != FREE) {
        waiting.enqueue(self);
        Enable_Interrupts();
        block();
        Disable_Interrupts();
    }
    status = BUSY;
    Enable_Interrupts();
}
```

```
class LOCK{
    int status = FREE;
    queue waiting;
}

LOCK::Release() {
    Disable_Interrupts();
    if (! waiting.empty()) {
        proc = waiting.dequeue();
        readyList.insert(proc);
    }
    status = FREE;
    Enable_Interrupts();
}
```

- `block()` is internal kernel call (kernel mode)
- Interrupts must be enabled before calling `block()`
- Similar to context switch
 - Calling process taken off CPU
 - Calling process stored in PCB
 - Calling process **NOT** put on ready list
- Better be on some other list before calling!

```
#include <pthread.h>

pthread_mutex_t lock;

pthread_mutexattr_t mattr;

pthread_mutexattr_init( &mattr );

pthread_mutex_init( &lock, &mattr );

pthread_mutex_lock( &lock );

pthread_mutex_unlock( &lock );

pthread_mutex_destroy( &lock );
```



```
pthread_mutex_t lockCounter;

int main ( void ) {
    ...
    pthread_mutex_init( &lockCounter, NULL );

    // launch threads
    ...
    // join threads
    ...
    pthread_mutex_destroy( & lockCounter);
    ...
    return 0;
}
```

```
void* thread_func ( void *param ) {  
    ...  
    pthread_mutex_lock(&lockCounter);  
    // execute CS()  
    ...  
    pthread_mutex_unlock(&lockCounter);  
    ...  
    pthread_exit(NULL);  
}
```

```
#include <mutex>
```

```
std::mutex lock;
```

```
// lock and unlock a section of code
```

```
lock.lock();
```

```
CS();
```

```
lock.unlock();
```

```
std::mutex lockCounter;  
  
int main ( void ) {  
    ...  
    // launch threads  
    ...  
    // join threads  
    ...  
    return 0;  
}
```

```
void thread_func ( ) {  
    ...  
    lockCounter.lock();  
    // execute CS()  
    lockCounter.unlock();  
    ...  
}
```

- Read/Modify/Write instruction
 - test-n-set instruction (atomic instruction)
- Functional Definition (really single instruction)

```
int testNset(int &var) {  
    int tmp;  
    tmp = var;  
    var = 1;  
    return tmp;  
}
```

Example:

1) inputVal = 0;

2) returnVal = testNset(inputVal);

Q1: What is the value of *inputVal* and *returnVal* after the execution of *testNset()*?

A1: inputVal = 1 and returnVal = 0.

Q2: What changes if *testNset()* is called again with the same value for *inputVal*?

- Every modern CPU has equivalent
 - exchange
 - compare-and-swap

```
LOCK::Acquire() {  
    // spins a process as long as  
    // status is true  
    while (testNset(status) == 1);  
}
```

“Spinlock”

```
LOCK::Release() {  
    status = 0;  
}
```

- Locks are good for MUTEX only
- General concept is needed to solve synchronization problems:



Semaphores

- Synchronization concept first proposed by Edsger Dijkstra.
- Semaphores are objects maintained by the OS.
- Each semaphore has a value and two atomic operations:
 - wait
 - signal

“test”

```
Proberen() { P() Wait() sem_wait()  
    while (count == 0);  
    count--;  
}
```

“increment”

```
Verhogen() { V() Signal() sem_post()  
    count++;  
}
```

```
int sem_wait( sem ) {  
    Disable_Interrupts();  
    sem.count--;  
    if (sem.count < 0) {  
        waitlist.enqueue(self);  
        Enable_Interrupts();  
        block();  
        Disable_Interrupts();  
    }  
    Enable_Interrupts();  
    return 0;  
}
```

```
int sem_post( sem ) {  
    Disable_Interrupts();  
    sem.count++;  
    if (sem.count <= 0) {  
        proc = waitlist.dequeue();  
        readyList.insert(proc);  
    }  
    Enable_Interrupts();  
    return 0;  
}
```

- Monitors are a high-level synchronization constructs.
- Monitors allow safe sharing of abstract data types among concurrent processes.
- Monitors provide synchronized methods that can only be entered by one process or thread at a time.

```
/**
 * Models a bank account with a balance.
 */
public class BankAccount {

    private double balance;

    /**
     * Constructs a bank account object with 0 balance.
     */
    public BankAccount() {    balance = 0.0;    }

    /**
     * Deposits a specified amount into the account.
     * @param amount - the amount of money to be deposited
     */
    public synchronized void deposit(double amount) {
        balance = balance + amount;
    }
}
```

- Monitors provide condition variables with two operations
 - wait: puts process or thread asleep
 - signal: triggers a waiting process or thread to resume operation exactly at the position where process or thread called wait

Condition Variables Implementation (pseudo code)


```
Wait( lock ) {  
    // put one thread asleep  
    lock.Acquire();  
}
```

```
Signal( lock ) {  
    // wake one sleeping thread  
    lock.release();  
}
```

```
Broadcast( lock ) {  
    // wake-up all n threads repeat n times  
    lock.release();  
}
```

Producer/Consumer Synchr. Problem

- Producer makes widgets, Consumer eats widgets
- Storage facility (limited size) for widgets
 - Storage is a circular data structure
- Producer must stop generating data when buffer is full.
- Consumer must stop reading data when no data are available.
- Examples: Networks, Disk I/O, etc.




```
sem_t fullB;  
sem_t emptyB;  
pthread_mutex_t mutex;  
  
void * producer(void *id) {  
    int myID = (int)id;  
    while(1) {  
        sem_wait( &emptyB );  
        pthread_mutex_lock( &mutex );  
        addWidgetToBuffer();  
        pthread_mutex_unlock( &mutex );  
        sem_post( &fullB );  
    }  
}
```

Q: What is the initial value of *emptyB* and *fullB*?

```
sem_t fullB;  
sem_t emptyB;  
pthread_mutex_t mutex;  
void * consumer(void *id) {  
    int myID = (int)id;  
    while(1) {  
        sem_wait( &fullB );  
        pthread_mutex_lock( &mutex );  
        eatWidgetFromBuffer();  
        pthread_mutex_unlock( &mutex );  
        sem_post( &emptyB );  
    }  
}
```

A: Initially, the value *emptyB* must be the size of the buffer and *fullB* must be 0.

```
sem_wait(&printer);  
sem_wait(&file);  
printFile();  
sem_post(&file);  
sem_post(&printer);
```

```
sem_wait(&file);  
sem_wait(&printer);  
printFile();  
sem_post(&printer);  
sem_post(&file);
```

- Deadlock is possible
- Be careful with order


- Small town
- One barber shop w/ one barber
- Barber always at work
- Sleeps when possible
- 1st customer in shop wakes barber
- Other customers wait
- All customers done, barber goes back to sleep

Sleeping Barber

```
Barber() {  
    while (1) {  
        barber.Wait();  
        mutex.Acquire();  
        numWaiting--;  
        mutex.Release();  
        done.Signal();  
    }  
}
```


```
Customer() {  
    mutex.Acquire();  
    if (numWaiting < CHAIRS) {  
        numWaiting++;  
        mutex.Release();  
        barber.Signal();  
        done.Wait();  
    } else  
        mutex.Release();  
    leave();  
}
```

Dining Philosophers (DP)



- N philosophers at round table
- Bowl of rice in middle of table
- Fork between each pair of philosophers
- Need 2 forks to eat
- think, get hungry, grab forks, eat, ...
- Avoid deadlock
- Avoid starvation
- Let as many eat as possible

Solution Outline




- Philosophers are identified by number
- Forks are identified by number
- P_i can grab F_i and $F_{i-1} \pmod N$
- Forks are non-preemptive resources.

How about this?

```
void *philosopherThread(void *id) {  
    int myID = (int)id;  
    while(1) {  
        think();  
        grab(myID - 1);  
        grab(myID);  
        eat();  
        replace(myID - 1);  
        replace(myID);  
    }  
}
```


Solution? Outline



- $P_0: G(F_0), G(F_4)$
- $P_1: G(F_1), G(F_0)$
- $P_3: G(F_3), G(F_2)$
- $P_4: G(F_4)$

- One philosopher's stomach growls
- All philosophers realize they are hungry
- All philosophers grab left fork
- All philosophers sleep waiting for right fork
- Deadlock!
- How can we fix the problem?


Solution? Outline



- $P_0: G(F_0), G(F_4)$
- $P_1: G(F_1), G(F_0)$
- $P_3: G(F_3), G(F_2)$
- $P_4: G(F_4), G(F_3)$
- $P_2: G(F_2), G(F_1)$

```
void *philosopherThread(void *id) {  
    int myID = (int)id;  
    while(1) {  
        think();  
        pthread_mutex_lock( &mutex );  
        grab(myID - 1);  
        grab(myID);  
        eat();  
        replace(myID - 1);  
        replace(myID);  
        pthread_mutex_unlock( &mutex );  
    }  
}
```

Solution? Outline



• $P_0: G(F_0), G(F_4)$

• $P_1: G(F_1)$

• $P_3: G(F_3)$

• $P_4: G(F_4)$

- Do we avoid deadlock?
- How many can eat at a time?
- Do you think this is a good solution?
- How can we improve it?
 - Analyze the problem
 - What is *really* causing deadlock
 - How can we avoid it?
 - Can we make our solution a *little* better?

Next Programming Assignment

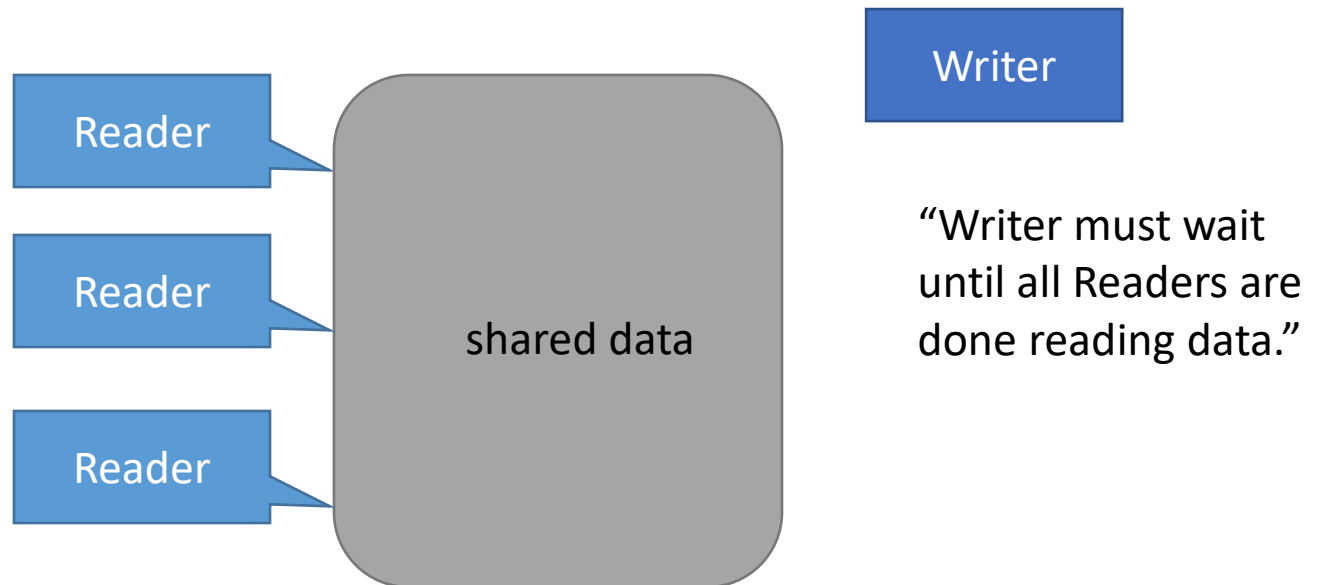



```
LOCK lock = FREE;
AddToQueue( item ) {
    lock.Acquire();
    list.append(item);
    lock.Release();
}
RemoveFromQueue() {
    lock.Acquire();
    if( ! list.isEmpty() )
        item = list.remove();
    lock.Release();
    return item;
}
```

```
RemoveFromQueue() {  
    lock.Acquire();  
    if( list.isEmpty() )  
        lock.Release();  
    Sleep();  
    lock.Acquire();  
    item = list.remove();  
    lock.Release();  
    return item;  
}
```

Multiple Reader, Single Writer

- Multiple Reader processes may access simultaneously shared data, provided Writer process has no access to it.
- Writer process may not access shared data unless no Reader process accesses shared data.



Example – DB Access (pseudo code)

```
monitor DB {  
    int activeReaders = 0;  
    int activeWriters = 0;  
    int waitingReaders = 0;  
    int waitingWriters = 0;  
    Condition OKtoRead;  
    Condition OKtoWrite;  
    Lock lock = FREE;
```

Example – DB Access (pseudo code)

```
ReadStart() {  
    lock.Acquire();  
    while((activeWriters + waitingWriters) > 0) {  
        waitingReaders++;  
        OKtoRead.Wait(lock);  
        waitingReaders--;  
    }  
    activeReaders++;  
    lock.Release();  
}
```

```
ReadStop() {  
    lock.Acquire();  
    activeReaders--;  
    if ((activeReaders == 0) && (waitingWriters > 0))  
        OKtoWrite.Signal(lock);  
    lock.Release();  
}
```

Example – DB Access (pseudo code)

```
WriteStart() {  
    lock.Acquire();  
    while((activeWriters + activeReaders) > 0) {  
        waitingWriters++;  
        OKtoWrite.Wait(lock);  
        waitingWriters--;  
    }  
    activeWriters++;  
    lock.Release();  
}
```

```
WriteStop() {  
    lock.Acquire();  
    activeWriters--;  
    if ( waitingWriters > 0 )  
        OKtoWrite.Signal(lock);  
    else if ( waitingReaders > 0 )  
        OKtoRead.Broadcast(lock);  
    lock.Release();  
}
```



```
Reader() {  
    while(1) {  
        ReadStart();  
        ReadDatabase();  
        ReadEnd();  
    }  
}
```

```
Writer() {  
    while(1) {  
        WriteStart();  
        WriteDatabase();  
        WriteEnd();  
    }  
}
```

Semaphore

- Has a value
- Implies a history
- `wait()` may sleep or may proceed depending on value
- `signal()` will increment a value and wake a sleeper
- `signal()` before `wait()` will effect result of wait

Condition Variable

- Doesn't have a value
- Implies no memory
- `wait()` always sleeps
- `signal()` will wake a sleeper
- `signal()` before `wait()` has no effect

- Threads/processes need to be synchronized to
 - prevent race conditions
 - establish an order of execution
- Monitors combined with condition variables and semaphores can be used as synchronization instruments.
- Many problems require thread and process synchronization:
 - read/writer problem
 - bounded buffer
 - mutual exclusion